

# Owicki-Gries Reasoning for C11 RAR

Anonymous

Anonymous

## Abstract

Owicki-Gries reasoning for concurrent programs uses Hoare logic together with an *interference freedom* rule for concurrency. In this paper, we develop a new proof calculus for the C11 RAR memory model (a fragment of C11 with both relaxed and release-acquire accesses) that allows all Owicki-Gries proof rules for compound statements, including non-interference, to remain unchanged. Our proof method features novel assertions specifying *thread-specific views* on the state of programs. This is combined with a set of Hoare logic rules that describe how these assertions are affected by atomic program steps. We demonstrate the utility of our proof calculus by verifying a number of standard C11 litmus tests and Peterson’s algorithm adapted for C11. Our proof calculus and its application to program verification have been fully mechanised in the theorem prover Isabelle.

**2012 ACM Subject Classification** To be done

**Keywords and phrases** C11, Verification, Hoare logic, Owicki-Gries, Isabelle

**Digital Object Identifier** 10.4230/LIPIcs.AAA.2020.0

## 1 Introduction

In 1976, Susan Owicki and David Gries proposed an extension of Hoare’s axiomatic reasoning technique [15] to concurrent programs [27]. Their proof calculus allows one to reason about concurrent programs with shared variables via a number of proof rules, including the rules for sequential programs as introduced by Hoare plus an additional proof rule for concurrent composition. This composition rule basically allows for the conjunction of pre- and post-conditions of the process’ individual proofs, given that their proof outlines are *interference free*. Interference freedom requires that an assertion in the proof of one process cannot be invalidated by a statement in another process, when executed under the statement’s precondition.

Today, concurrent programs are run on multi-core processors. Multi-core processors come with *weak memory models* specifying the execution behaviour of concurrent programs. Reasoning consequently needs to be adapted to the memory model under consideration. Owicki-Gries reasoning is, however, fixed to the memory model of *sequential consistency* (SC) [23], and is unsound for weak memory models. Recent research has thus worked towards new sound proof calculi for concurrent programs. Most often, such approaches involve concurrent separation logics (e.g., GPS and RSL [34, 16]). These techniques constitute a radical departure from the (relatively) small and easy proof calculus of Owicki and Gries, further extending already complex logics. A proposal for a (rely-guarantee variant of) the Owicki-Gries proof system has been made by Lahav and Vafeiadis [22], however, requiring a strengthened non-interference check.

In this paper, we develop a proof method based on the Owicki-Gries proof calculus, keeping all of the original proof rules including the non-interference check unchanged. Our technique introduces a set of basic axioms to cope with memory accesses (reads, writes, read-modify-writes) and simple assertions that describe the current configuration of the weak memory state. Our proof calculus targets the weak memory model of the C11 programming language [9]. Here, we deal with the release-acquire-relaxed (RAR) fragment of C11 (thereby going further than prior work on Owicki-Gries reasoning for C11 [22]).

The key idea of our approach is the usage of novel assertions, which allow to specify



© Anonymous;  
licensed under Creative Commons License CC-BY

AAA.

Editors: AAA; Article No. 0; pp. 0:1–0:38



Leibniz International Proceedings in Informatics

Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

*thread-specific* views on shared variables. We also include a specific assertion containing a modality for release-acquire (RA) synchronisation, capturing particularities of C11 RA message passing. The use of non-standard assertions as a consequence necessitates the introduction of new rules of assignment, formalising the effect of assignments on assertions.

We build our proof calculus on top of an *operational* semantics for C11 RAR. The semantics is a mixture of the operational semantics proposed by Doherty et al. [12] (for RAR) and Kaiser et al.'s semantics [16] for RA plus non-atomics. Correctness of this novel proposal is shown by proving it to coincide with the semantics in [12] which in turn has been proven to coincide with the standard axiomatic semantics of Batty et al. [9]. We have formalised our semantics within the theorem prover Isabelle [28] and mechanically proved soundness of all of our new rules for C11 assertions. Moreover, we provide mechanical proofs<sup>1</sup> of several litmus tests from the literature (message passing, load buffering, read-read coherence) as well as a version of Peterson's algorithm adapted for C11 memory [12, 36].

*Overview.* The paper is organised as follows. In the next section we start with an example explaining the behaviour of concurrent programs on C11, motivating our novel assertions. Section 3 defines the syntax of C11 RAR programs and Section 4 its semantics. We present the proof calculus and its novel assertions in Section 5 via proofs of correctness for some standard litmus tests, and a case study of Peterson's algorithm in Section 6. Section 7 describes our Isabelle mechanisation, Section 8 discusses related work and the last section concludes.

## 2 Deductive Reasoning for Weak Memory

In this section, we illustrate the basic principles of C11 synchronisation and our verification method by considering the message-passing example (Figures 1 and 2). The two programs are almost identical and consist of two threads executing in parallel, accessing shared variables. The assertions in curly brackets at the end specify the programs' postconditions.

The programs comprise two shared variables:  $d$  (that stores some data) and  $f$  (that stores a flag). In both programs, both  $d$  and  $f$  are initially 0. thread 1 updates  $d$  to 5, then updates  $f$  to 1. Thread 2 waits for  $f$  to be set to 1, then reads from  $d$ . Under sequential consistency, one would expect that the final value of  $r2$  is 5, since the loop in thread 2 only terminates after  $f$  has been updated to 1 in thread 1, which in turn happens after  $d$  has been set to 5. However, the C11 semantics allows the behaviour in Figure 2, where thread 2 may read a stale value of  $d$ , and hence only the weaker postcondition  $r2 = 0 \vee r2 = 5$  holds. To regain the expected behaviour, one must introduce additional synchronisation in the program. In particular, the write to  $f$  by thread 1 must be a *releasing write* (i.e.,  $f :=^R 1$ ) and the read of  $f$  in thread 2 must be an *acquiring read* (i.e.,  $r1 \leftarrow^A f$ ) as in Figure 1.

In sequential consistency all threads have a single common view of the shared state, namely all threads see the latest write that occurs for each variable. When a new write is executed, the views of all threads are updated so that they see this write. In contrast, each thread in C11 programs has its own view of each variable, which is affected by synchronisation annotations. Thus, for the program in Figure 2, after initialisation, all threads see the initial writes (i.e.,  $d = 0, f = 0$ ). The assignments in thread 1 only change thread 1's view, and leave thread 2's view unchanged. Thus, after execution of  $f := 1$ , thread 2 has access to two

<sup>1</sup> The Isabelle files may be downloaded from: <https://www.dropbox.com/sh/4yr2w7792qwyw09/AACsWUXtZbK3PvyfJkqjyDYa> within the file ECOOP-2020-Isabelle.zip.

**Init:**  $d := 0; f := 0;$   
**Thread 1**  $\parallel$  **Thread 2**  
 $d := 5;$   $\parallel$  **do**  $r1 \leftarrow^A f$   
 $f :=^R 1;$   $\parallel$  **until**  $r1 = 1;$   
 $\parallel$   $r2 \leftarrow d;$   
 $\{r2 = 5\}$

■ **Figure 1** Message-passing litmus test

**Init:**  $d := 0; f := 0;$   
**Thread 1**  $\parallel$  **Thread 2**  
 $d := 5;$   $\parallel$  **do**  $r1 \leftarrow f$   
 $f := 1;$   $\parallel$  **until**  $r1 = 1;$   
 $\parallel$   $r2 \leftarrow d;$   
 $\{r2 = 0 \vee r2 = 5\}$

■ **Figure 2** Unsynchronised message passing

88 values for  $d$  (i.e.,  $d \in \{0, 5\}$ ) and  $f$  (i.e.,  $f \in \{0, 1\}$ ). Even if thread 2 reads  $f = 1$ , its view of  
 89  $d$  remains unchanged and it continues to have access to both values of  $d$ .

90 The program in Figure 1 has a similar semantics for initialisation and execution of thread 1,  
 91 i.e., its execution does not affect the view of thread 2. However, due to the release-acquire  
 92 synchronisation on  $f$  (notation R and A), after thread 2 reads  $f = 1$ , its view for  $d$  will be  
 93 updated so that the stale value  $d = 0$  is no longer available for it to read. One way to explain  
 94 this behaviour is by thinking of thread 1 as *passing its knowledge of the write* to  $d$  to thread  
 95 2 via the variable  $f$ , which is synchronised using the release-acquire annotations.

96 This intuition is captured formally using a semantics based on *timestamps* [16, 13, 17, 29],  
 97 which enables one to encode each thread's view and define how these views are updated. In  
 98 this paper, we characterise the release-acquire-relaxed subset of C11 [12] (C11 RAR) using  
 99 timestamps, which has a restriction prohibiting the so-called *load-buffering* litmus test [20].

100 The main contribution of our paper is an assertion language that enables one to reason  
 101 about thread views in a Hoare-style proof calculus, resulting in the proof outline given in  
 102 Figure 3. As already noted, the key advantage of these assertions is the fact that standard  
 103 rules of Hoare and Owicki-Gries logic remain unchanged. For message passing, we require  
 104 three main types of assertions (see Section 5):

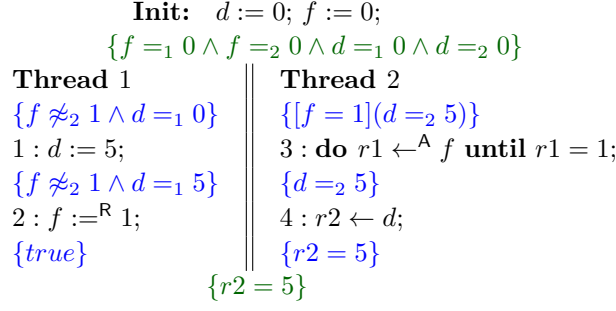
105 **Possible value.** A possible value assertion (denoted  $x \approx_t n$ ) states that thread  $t$  can read  
 106 value  $n$  of global variable  $x$ , i.e., there is a write to  $x$  with value  $n$  beyond or including  
 107 the *viewfront*<sup>2</sup> of thread  $t$ . Note that there may be more than one such write, and hence  
 108 there may be several possible values for a given variable. Moreover, the last write to each  
 109 variable is always viewable as a possible value.

110 **Definite value.** A definite value assertion (denoted  $x =_t n$ ) states that thread  $t$ 's viewfront  
 111 is up-to-date with the writes to  $x$  (i.e., there is a single write to  $x$  beyond or including  
 112 the viewfront of thread  $t$ ), and this write updates  $x$ 's value to  $n$ . Thus,  $t$  definitely knows  
 113 the variable  $x$  to have value  $n$ .

114 **Conditional value.** A conditional value assertion (denoted  $[x = n](y =_t m)$ ) captures the  
 115 message passing idiom for variable  $y$  via variable  $x$ . It guarantees that when thread  $t$   
 116 reads  $x$  to be  $n$  via an acquiring read, a release-acquire synchronisation is induced and  
 117 thereby  $t$  learns the definite value of  $y$  to be  $m$ . In particular, after reading  $x = n$  via  
 118 an acquiring read, the viewfront for  $t$  is updated so that the only write to  $y$  beyond or  
 119 including this viewfront is a write with value  $m$ .

120 For the example in Figure 3, after initialisation, both threads 1 and 2 have definite value 0  
 121 for both  $d$  and  $f$ . The precondition of  $d := 5$  states that thread 2 cannot possibly observe 1  
 122 for  $f$  (i.e.,  $f \approx_2 1$ ) and thread 1 definitely observes 0 for  $d$  (i.e.,  $d =_1 0$ ). These assertions can  
 123 be proven *locally correct* and *interference free* since thread 2 neither modifies  $d$  nor  $f$ . The

<sup>2</sup> We borrow the term viewfront from Popkadaev et al. [29].



■ **Figure 3** Proof outline for message passing

precondition of  $f :=^R 1$  is similar but with  $d =_1 5$  in place of  $d =_1 0$ . The precondition of the **until** loop in thread 2 contains a conditional value assertion, which ensures that if thread 2 reads  $f = 1$  then it will definitely read  $d = 5$ . This conditional value assertion enables one to establish local correctness of the precondition (i.e.,  $d =_2 5$ ) of the statement  $r2 \leftarrow d$ , which leads to the postcondition of the program. Each of the assertions in thread 2 can be proven to be interference free against thread 1.

### 3 Program Syntax

We start by defining the syntax of concurrent programs, starting with the structure of sequential programs (single threads). A thread may use *global* shared variables (from  $Var_G$ ) and local registers (from  $Var_L$ ). We let  $Var = Var_G \cup Var_L$  and assume  $Var_G \cap Var_L = \emptyset$ . Global variables can be accessed in three different *synchronisation modes*: acquire (A, for reads), release (R, for writes) and relaxed (no annotation). The annotation RA is employed for *update* operations, which read and write to a shared variable in a single atomic step. We use  $x, y, z$  to range over global variables and  $r1, r2, \dots$  to range over local variables. We assume that  $\ominus$  is a unary operator (e.g.,  $\neg$ ),  $\oplus$  is a binary operator (e.g.,  $\wedge, +, =$ ) and  $n$  is a value (of type  $Val$ ). Expressions may only involve local variables. For a treatment of expressions with global variables in the semantics see [12]. The syntax of sequential programs,  $Com$ , is given by the following grammar (with  $r \in Var_L, x \in Var_G$ ):

$$\begin{aligned}
 Exp_L &::= Val \mid r \mid \ominus Exp_L \mid Exp_L \oplus Exp_L \\
 ACom &::= \text{skip} \mid x.\text{swap}(n)^{RA} \mid r := Exp_L \mid x :=^{[R]} Exp_L \mid r \leftarrow^{[A]} x \\
 Com &::= ACom \mid Com; Com \mid \text{if } B \text{ then } Com \text{ else } Com \mid \text{while } B \text{ do } Com
 \end{aligned}$$

where we assume  $B$  to be an expression of type  $Exp_L$  that evaluates to a boolean. The statement  $x.\text{swap}(n)^{RA}$  atomically reads the variable  $x$  (using an acquiring read) and updates  $x$  to value  $n$  (using a releasing write) in a single atomic step. Its execution therefore gives rise to an atomic read-modify-write update event. We have not included a **CAS** operation here; it could similarly be implemented by an update event (see e.g. [35]).

The notation  $[X]$  denotes that the annotation  $X$  is optional, where  $X \in \{A, R\}$ , enabling one to distinguish relaxed, acquiring and releasing accesses. Loops will be used in other forms, like **do-until** or **do-while**, which are straightforward to define in terms of the command syntax above.

As is standard in Owicki-Gries proofs, we make use of *auxiliary variables*, which are variables that do not affect the meaning of a program, but appear in proof assertions. We require that each auxiliary variable is *local* to the thread in which it occurs. Auxiliary

variables may only occur in assignments, not in conditional statements, and only in the form  $a := E$ , where  $E \in \text{Exp}_L$  and  $a$  is an auxiliary variable<sup>3</sup>. Finally, we require that writes to auxiliary variables occur atomically in conjunction with another (non-auxiliary) atomic program step. Such atomic operations are written as  $\langle A, a := E \rangle$ , where  $A \in \text{ACom}$ . This is more of a technical requirement which could also easily be relaxed. It guarantees that the programs without and with auxiliary variables have the same number of transitions (no stuttering steps).

For simplicity, we assume concurrency at the top level only. We let  $\text{Tid}$  be the set of all thread identifiers and use a function  $\text{Prog} : \text{Tid} \rightarrow \text{Com}$  to model a program comprising multiple threads. In examples, we typically write concurrent programs as  $C_1 \parallel \dots \parallel C_n$ , where  $C_i \in \text{Com}$ . We further assume some initialisation of variables. The structure of our programs thus is **Init**;  $(C_1 \parallel \dots \parallel C_n)$ .

## 4 Semantics

The operational semantics for this language is defined in two parts. The *program semantics* fixes the steps that the concurrent program can take. This gives rise to transitions  $(P, \text{lst}) \xrightarrow{a}_t (P', \text{lst}')$  of a thread  $t$  where  $P$  and  $P'$  are programs,  $\text{lst}$  and  $\text{lst}'$  is the state of local variables and  $a$  is an action (possibly the silent action  $\tau$ , see below). The program semantics is combined with a *memory semantics* which reflects the C11 state (denoted by  $\sigma$ ), and in particular the write actions from which a read action can read.

We start by fixing the actions, where  $x \in \text{Var}_G$  and  $m, n \in \text{Val}$ :

$$\text{Act} = \{rd(x, n), rd^A(x, n), wr(x, n), wr^R(x, n), upd^{RA}(x, n, m)\}$$

containing actions for (releasing) reads, (acquiring) writes and updates (reading value  $n$  and writing  $m$ ). We furthermore employ a silent  $\tau$  action and let  $\text{Act}_\tau = \text{Act} \cup \{\tau\}$ . For an action  $a \in \text{Act}$ , we let  $\text{var}(a) \in \text{Var}_G$  be the variable read (or written to),  $\text{rdval}(a) \in \text{Val}$  be the value read and  $\text{wrval}(a) \in \text{Val}$  be the value written. We let  $\text{U}$  denote the update actions, and distinguish the sets  $\text{W}_R \supseteq \text{U}$  (write release),  $\text{R}_A \supseteq \text{U}$  (read acquire),  $\text{W}_X$  (write relaxed) and  $\text{R}_X$  (read relaxed). Finally, we define  $\text{R} = \text{R}_A \cup \text{R}_X$  (all reads) and  $\text{W} = \text{W}_R \cup \text{W}_X$  (all writes). Typically, we refer to the elements of  $\text{W}$  as *writes*, but note that this set also includes update actions.

### 4.1 Program Semantics

In the program semantics, we assume a function  $\text{lst} \in \text{Tid} \rightarrow (\text{Var}_L \leftrightarrow \text{Val})$ , which returns the local state for the given thread. We assume that the local variables of threads are disjoint, i.e., if  $t \neq t'$ , then  $\text{dom}(\text{lst}(t)) \cap \text{dom}(\text{lst}(t')) = \emptyset$ . For an expression  $E$  over local variables, we write  $\llbracket E \rrbracket_{ls}$  for the value of  $E$  in local state  $ls$ ; we write  $ls[r := n]$  to state that  $ls$  remains unchanged except for the value of local variable  $r$  which becomes  $n$ .

Figure 4 gives the transition rules of the program semantics. The last rule, **Prog**, lifts the transitions of threads to a transition for a concurrent program. The other rules concern the sequential part of the language. The rules in a sense ignore the fact that the language allows for global variables; the program semantics just details the values of local variables in component  $ls$ . When global variables are read, the program semantics allows for *all* possible

<sup>3</sup> The locality requirement is the only difference to “normal” Owicki-Gries auxiliary variables.

$$\begin{array}{c}
\frac{r \in \text{Var}_L \quad n = \llbracket E \rrbracket_{ls}}{(r := E, ls) \xrightarrow{\tau} (\mathbf{skip}, ls[r := n])} \quad \frac{x \in \text{Var}_G \quad a = wr^{[R]}(x, \llbracket E \rrbracket_{ls})}{(x :=^{[R]} E, ls) \xrightarrow{a} (\mathbf{skip}, ls)} \\
\\
\frac{a = rd^{[A]}(x, n) \quad n \in \text{Val}}{(r \leftarrow^{[A]} x, ls) \xrightarrow{a} (\mathbf{skip}, ls[r := n])} \quad \frac{a = upd^{\text{RA}}(x, m, n) \quad m \in \text{Val}}{(x.\mathbf{swap}(n)^{\text{RA}}, ls) \xrightarrow{a} (\mathbf{skip}, ls)} \\
\\
\frac{(C_1, ls) \xrightarrow{a} (C'_1, ls')}{(C_1; C_2, ls) \xrightarrow{a} (C'_1; C_2, ls')} \quad \frac{}{(\mathbf{skip}; C_2, ls) \xrightarrow{\tau} (C_2, ls)} \\
\\
\frac{\llbracket B \rrbracket_{ls}}{(\mathbf{if } B \mathbf{ then } C_1 \mathbf{ else } C_2, ls) \xrightarrow{\tau} (C_1, ls)} \quad \frac{\neg \llbracket B \rrbracket_{ls}}{(\mathbf{if } B \mathbf{ then } C_1 \mathbf{ else } C_2, ls) \xrightarrow{\tau} (C_2, ls)} \\
\\
\frac{\llbracket B \rrbracket_{ls}}{(\mathbf{while } B \mathbf{ do } C, ls) \xrightarrow{\tau} (C; \mathbf{while } B \mathbf{ do } C, ls)} \quad \frac{\neg \llbracket B \rrbracket_{ls}}{(\mathbf{while } B \mathbf{ do } C, ls) \xrightarrow{\tau} (\mathbf{skip}, ls)} \\
\\
\text{AUX} \frac{(A, ls) \xrightarrow{a} (\mathbf{skip}, ls') \quad (a := E, ls') \xrightarrow{\tau} (\mathbf{skip}, ls'')}{(\langle A; a := E \rangle, ls) \xrightarrow{a} (\mathbf{skip}, ls'')} \quad \text{PROG} \frac{(P(t), lst(t)) \xrightarrow{a} (C, ls) \quad a \in \text{Act}_\tau}{(P, lst) \xrightarrow{a}_t (P[t := C], lst[t := ls])}
\end{array}$$

■ **Figure 4** Program semantics

values to be read. This is combined with the memory semantics (formalised by  $\xrightarrow{a}_t$ ) as follows:

$$\frac{(P, lst) \xrightarrow{\tau}_t (P', lst')}{(P, lst, \sigma) \Longrightarrow (P', lst', \sigma)} \quad \frac{(P, lst) \xrightarrow{a}_t (P', lst') \quad \sigma \xrightarrow{a}_t \sigma'}{(P, lst, \sigma) \Longrightarrow (P', lst', \sigma')}$$

The transitions defined by  $\sigma \xrightarrow{a}_t \sigma'$  ensure that read actions only return a value allowed by the C11 semantics and are defined in Section 4.2. The rules for all imperative program constructs (sequential composition, **if** and **while**) are standard.

## 4.2 Memory Semantics

Next, we detail the memory semantics, which is equivalent to an earlier operational reformulation [12] of the RAR fragment from [20].

**C11 State.** Table 1 summarises the components of a C11 state. Each global write is represented by a pair  $(a, q) \in \mathcal{W} \times \mathbb{Q}$ , where  $a$  is a write action, and  $q$  is a rational number that we use as a *timestamp* (c.f., [16, 13, 29]). The timestamps totally order the writes to each variable; the ordering induced by timestamps is also referred to as the *modification order* [20, 12] or *coherence order* [3]. For each write  $w = (a, q)$ , we denote  $w$ 's timestamp by  $tst(w) = q$ . We also lift the functions *var* and *wrval* to timestamped writes, e.g.,  $var((a, q)) = var(a)$ . The set of all writes that have occurred in the execution thus far is recorded in the state component  $writes \subseteq \mathcal{W} \times \mathbb{Q}$ .

As described in Section 2, each state must record the writes that are observable to each read. To achieve this, we use two families of functions from global variables to writes, both of which record the *viewfronts* (c.f., [29, 17]).

■ A function  $twiew_t$  that returns the *viewfront* of thread  $t$ . The thread  $t$  can read from any write to variable  $x$  whose timestamp is not earlier than  $twiew_t(x)$ . Accordingly, we define,

■ **Table 1** Components of a C11 state

Component	Informal meaning	Initial value
$writes \subseteq W \times \mathbb{Q}$	The writes which have happened so far	$writes_{\mathbf{Init}}$
$tview_t \in Var_G \rightarrow writes$	The view of a thread $t$	$tview_{\mathbf{Init}}$
$mview_w \in Var_G \rightarrow writes$	The view of a thread when writing $w$	$mview_{\mathbf{Init}}$
$covered \subseteq writes$	The covered writes	$\emptyset$

218 for each state  $\sigma$ , thread  $t$  and global variable  $x$ , the set of *observable writes*:

$$219 \quad \sigma.OW(t, x) = \{(a, q) \in \sigma.writes \mid var(a) = x \wedge tst(\sigma.tview_t(x)) \leq q\} \quad (1)$$

221 ■ A function  $mview_w$  that records the *viewfront* of write  $w$ , which is set to be the viewfront  
 222 of the thread that executed  $w$  at the time of  $w$ 's execution. We use  $mview_w$  to compute  
 223 a new value for  $tview_t$  if a thread  $t$  *synchronizes* with  $w$ , i.e., if  $w \in W_R$  and another  
 224 thread executes an  $e \in R_A$  that reads from  $w$ .

225 Finally, our semantics maintains a variable  $covered \subseteq writes$ . In C11 RAR, each update  
 226 action occurs in modification order immediately after the write that it reads from [12]. This  
 227 property constitutes the atomicity of updates. In order to preserve this property, we must  
 228 prevent any newer write from intervening between any update and the write that it reads  
 229 from. As we explain below, *covered* writes are those that are immediately prior to an update  
 230 in modification order, and new write actions never interact with a covered write.

231 **Initialisation.** Table 1 also states how these components are initialised by **Init**. If  $Var_G =$   
 232  $\{x_1, \dots, x_n\}$ ,  $Var_L = \{r_1, \dots, r_m\}$  and  $k_1, \dots, k_n, l_1, \dots, l_m \in Val$ , we assume **Init** =  $x_1 :=$   
 233  $k_1; \dots, x_n := k_n; [r_1 := l_1;] \dots [r_m := l_m;]$ , where we use the notation  $[r_i := l_i;]$  to mean  
 234 that the assignment  $r_i := l_i$  may optionally appear in **Init**. Thus each shared variable is  
 235 initialised exactly once and each local variable is initialised at most once. The initial values  
 236 of the state components are then as follows, where we assume that 0 is the initial timestamp.

$$237 \quad writes_{\mathbf{Init}} = \{(wr(x_1, k_1), 0), \dots, (wr(x_n, k_n), 0)\}$$

$$238 \quad tview_{\mathbf{Init}}(x_i) = (wr(x_i, k_i), 0) \quad \text{for each thread } x_i \in Var_G$$

$$239 \quad mview_{\mathbf{Init}} = tview_{\mathbf{Init}}$$

241 The local state component of each thread must also be compatible with **Init**, i.e., for each  $t$   
 242 if  $r_i \in \mathbf{dom}(lst(t))$  we have that  $(lst(t))(r_i) = l_i$  provided  $r_i := l_i$  appears in **Init**.

243 We let  $lst_{\mathbf{Init}}$  be the local state compatible with **Init**, let  $\sigma_{\mathbf{Init}}$  denote the initial state  
 244 defined by **Init**, and define  $\Gamma_{\mathbf{Init}} = (lst_{\mathbf{Init}}, \sigma_{\mathbf{Init}})$ .

245 **Transition semantics.** The transition relation of our semantics for global reads and writes  
 246 is given in Figure 5. Each transition  $\sigma \xrightarrow{a}_t \sigma'$  is labelled by an action  $a$  and thread  $t$ . The  
 247 premise of each rule must identify the write  $w$  that the action interacts with. This is made  
 248 more precise below.

249 **READ transition by thread  $t$ .** Here we assume that

- 250 ■  $a$  is either a relaxed or acquiring read to variable  $x$ ,
- 251 ■  $w$  is a write to  $x$  that  $t$  can observe (i.e.,  $(w, q) \in \sigma.OW(t, x)$ ), and
- 252 ■ the value read by  $a$  is the value written by  $w$ .

253 Each read causes the viewfront of  $t$  to be updated. This is computed as follows. If the read  
 254 synchronises with the write, then the thread's new view will be a combination of its existing



$$\begin{array}{c}
\text{READ} \frac{a \in \{rd(x, n), rd^A(x, n)\} \quad (w, q) \in \sigma.OW(t, x) \quad wrval(w) = n}{\sigma \xrightarrow{a}_t \sigma[tview_t := tview'_t]} \\
\text{WRITE} \frac{a \in \{wr(x, n), wr^R(x, n)\} \quad (w, q) \in \sigma.OW(t, x) \setminus \sigma.covered \quad \sigma.fresh(q, q') \quad \begin{array}{l} tview'_t = \begin{cases} \sigma.tview_t \otimes \sigma.mview_{(w, q)} & \text{if } (w, a) \in W_R \times R_A \\ \sigma.tview_t[x := (w, q)] & \text{otherwise} \end{cases} \\ writes' = \sigma.writes \cup \{(a, q')\} \quad tview'_t = \sigma.tview_t[x := (a, q')] \end{array}}{\sigma \xrightarrow{a}_t \sigma[tview_t := tview'_t, mview_{(a, q')} := tview'_t, writes := writes']} \\
\text{UPDATE} \frac{a = upd^{RA}(x, m, n) \quad (w, q) \in \sigma.OW(t, x) \setminus \sigma.covered \quad wrval(w) = m \quad \sigma.fresh(q, q') \quad \begin{array}{l} writes' = \sigma.writes \cup \{(a, q')\} \quad covered' = \sigma.covered \cup \{(w, q)\} \\ tview'_t = \begin{cases} \sigma.tview_t[x := (a, q')] \otimes \sigma.mview_{(w, q)} & \text{if } w \in W_R \\ \sigma.tview_t[x := (a, q')] & \text{otherwise} \end{cases} \end{array}}{\sigma \xrightarrow{a}_t \sigma[tview_t := tview'_t, mview_{(a, q')} := tview'_t, writes := writes', covered := covered']}
\end{array}$$

■ **Figure 5** Transition relation of the memory semantics

view, and the view of that write. In particular, for each variable  $x$  the new view of  $x$  will be the later of either  $tview_t(x)$  or  $mview_w(x)$ , in timestamp order. To express this, we use an operation that combines two views  $v_1$  and  $v_2$ , by constructing a new view that takes the later of the writes at each variable:

$$(v_1 \otimes v_2)(x) = \begin{cases} v_1(x) & \text{if } tst(v_2(x)) \leq tst(v_1(x)) \\ v_2(x) & \text{otherwise} \end{cases}$$

If  $w$  and  $a$  do not synchronise, then  $tview_t$  is simply updated to include the new write.


For illustration, consider the picture in Figure 6. The x-axis depicts the timestamps of the writes, the y-axis the variables  $x, y$  and  $z$ , which we assume are initialised by writes  $x_0, y_0$  and  $z_0$ , respectively. The orange line shows the view of a thread, say  $t_1$ , and the blue line depicts the view of another thread that executes  $w = (wr^R(y, 42), 3)$ . If thread  $t_1$  performs an acquiring read of  $y$  and reads from  $w$  (i.e., it performs a synchronising read), thread  $t_1$ 's view changes to the diagram on the right, whereby its current viewfront is combined with the viewfront of  $w$ .

**WRITE transition by thread  $t$ .** A write transition must identify the write  $(w, q)$  after which  $a$  occurs. This  $w$  must be observable and must *not* be covered — the second condition is required to preserve the read-modify-write atomicity of updates. We must choose a fresh timestamp  $q' \in \mathbb{Q}$  for  $a$ , which is formalised by  $fresh(q, q')$ :

$$\sigma.fresh(q, q') = q < q' \wedge \forall w' \in \sigma.writes. q < tst(w') \Rightarrow q' < tst(w')$$

The predicate  $fresh(q, q')$  ensures that  $q'$  is a new timestamp for the variable  $x$ , such that  $(a, q')$  occurs immediately after  $(w, q)$ . The new write is added to the set  $writes$ . We update  $tview_t$  to include the new write, which means  $t$  can no longer observe any writes prior to  $(a, q')$ . Finally, we set the viewfront of  $(a, q')$  to be the new viewfront of  $t$ , i.e.,  $mview_{(a, q')} := tview'_t$ . Now, if some other thread synchronises with this new write in some later transition, that thread's view will become at least as recent as  $t$ 's view at this transition.





■ **Figure 6** Illustration of views and view updates: pre-state (left) and post-state (right)

280 **UPDATE transition by thread  $t$ .** These transitions are best understood as a combination  
 281 of the read and write transitions. As with a write transition, we must choose a valid fresh  
 282  $q'$ , and the state components *writes* and *mview* are updated in the same way. As discussed  
 283 earlier, in UPDATE transitions it is necessary to record that the write that the update interacts  
 284 with is now covered, which is achieved by adding that write to *covered*. Finally, we must  
 285 compute a new thread view, which is similar to a READ transition, except that the thread's  
 286 new view always includes the new write introduced by the update.

### 287 4.3 Relationship to the Axiomatic Semantics

288 We prove that the timestamp-based semantics presented here is equivalent to an earlier  
 289 operational semantics [12] that is already known to be equivalent to the C11 RAR fragment.  
 290 Here, we just roughly sketch how this proof proceeds, the appendix contains more details.

291 The semantics in [12] describes C11 states in the form  $E = (X, \text{sb}, \text{rf}, \text{mo})$ , where  $X$  is  
 292 a set of read and write events (roughly equivalent to actions) and *sb*, *rf* and *mo* describe  
 293 the sequenced-before and reads-from relation as well as the modification order of the C11  
 294 axiomatic semantics. A number of further relations are derived from these, in particular the  
 295 extended coherence order *eco* and the happens-before order *hb*. The proof of equivalence of  
 296 the semantics shows the two semantics to *simulate* each other. For this, we need to define a  
 297 correspondence between C11 states of form  $E$  and of form  $\sigma$  such that: (1) For  $\sigma.\text{writes}$ , we  
 298 take  $X \cap W$ ; (2) For  $\sigma.\text{covered}$ , we take the writes  $w$  in  $X \cap W$  such that there is an update  $u$   
 299 with  $(w, u) \in \text{rf}$ ; and (3) For *mview* and *tview*, we use a downward closure operator, **cclose**,  
 300 which for a given set of events  $S$  determines the set of events prior to  $S$  in the relation  $\text{eco}^? \circ \text{hb}^?$ .  
 301 Then  $\sigma.\text{tview}_t = \text{max}_{\text{mo}}(X.\text{cclose}(X_t))$  and  $\sigma.\text{mview}_w = \text{max}_{\text{mo}}(X.\text{cclose}(\{w\}))$ , where  
 302  $\text{max}_{\text{mo}}$  selects writes being maximal wrt. *mo* and  $X_t$  are all actions of  $t$  in  $X$ . In all these  
 303 cases, timestamps for writes have to be selected consistent with *mo*.

304 Given such a correspondence, the proof proceeds by showing this correspondence is  
 305 preserved by the read, write and update transitions.

### 306 4.4 Well Formedness

307 Our proofs in subsequent sections require that the state under consideration is *well-formed*.  
 308 This is formalised by predicate *wfs* over a C11 state  $\sigma$ , where

$$\begin{aligned}
 \text{wfs}(\sigma) \iff & \text{ran}((\bigcup_t \sigma.\text{tview}_t) \cup (\bigcup_w \sigma.\text{mview}_w)) \subseteq \sigma.\text{writes} \wedge \\
 & \text{finite}(\sigma.\text{writes}) \wedge \sigma.\text{covered} \subseteq \sigma.\text{writes} \wedge \\
 & (\forall w. w \in \sigma.\text{writes} \Rightarrow \sigma.\text{mview}_w(\text{var}(w)) = w)
 \end{aligned}$$

313 The first conjunct ensures that each viewable write is in  $\sigma.writes$ . The second conjunct  
 314 ensures there are only a finite number of writes, and the third ensures that every covered  
 315 write is an actual write. The final conjunct ensures that for each write in  $\sigma.writes$ , the  
 316 viewfront of  $w$  for  $var(w)$  is  $w$  itself.

317 Well-formedness is invariant for any program, i.e., every initialisation establishes well-  
 318 formedness and every program transition preserves well-formedness.

319 ► **Lemma 1.** *For any program  $C$  constructed using the syntax described in Section 3,  $wfs(\sigma)$*   
 320 *is invariant.*

321 **Proof.** In Isabelle. We show that every initialisation establishes  $wfs(\sigma)$ . Furthermore, if  
 322  $wfs(\sigma)$  and  $\sigma \xrightarrow{a}_t \sigma'$ , then  $wfs(\sigma')$  for any action  $a$  and thread  $t$ . ◀

## 323 5 Hoare Logic and Owicki-Gries Reasoning for C11

324 In this section, we present a Hoare logic [15] for C11 RAR that enables Owicki-Gries  
 325 reasoning [27]. For compound statements (including concurrent composition) we use the  
 326 standard rules of Hoare logic as well as the standard interference freedom proof obligations  
 327 described by Owicki and Gries. Our contribution is a novel set of high-level predicates  
 328 that describe the *observations* of each thread for a C11 state, together with a set of *basic*  
 329 *axioms* that describe how these predicates interact with read, write and update transitions.  
 330 Soundness of these axioms has been checked using Isabelle.

331 In Section 5.1, we link our operational semantics to the proof outlines of Hoare logic  
 332 and Owicki-Gries' notion of interference freedom. Section 5.2 provides an overview of our  
 333 assertion language and briefly discusses the main categories of assertions, i.e., assertions  
 334 describing observability, ordering and occurrences of writes. We present the basic axioms  
 335 in stages, using specific litmus tests (in Sections 5.3, 5.4, 5.5) to motivate each group of  
 336 assertions. The proof outlines of all litmus tests have been verified using Isabelle.

### 337 5.1 Soundness and Classical Verification Rules

338 We first define the meaning of a Hoare triple under partial correctness and present the  
 339 classical proofs rules for compound statements. Unlike Hoare logic, where a state is modelled  
 340 by a mapping from variables to values, as we have seen in Section 4.1, states of a C11  
 341 program contain two components: a local state  $lst$  and a global state  $\sigma$ . We let  $\Sigma_G$   
 342 be the set of all possible global state configurations (as described in Table 1) and let  
 343  $\Sigma_{C11} = (Var_L \rightarrow Val) \times \Sigma_G$  be the set of all possible C11 states. Predicates over  $\Sigma_{C11}$  are  
 344 therefore of type  $\Sigma_{C11} \rightarrow \mathbb{B}$ . This leads to the following definition of a Hoare triple, which we  
 345 note is the same as the standard definition — the only difference is that the state component  
 346 is of type  $\Sigma_{C11}$ .

347 ► **Definition 2.** *Suppose  $p, q \in \Sigma_{C11} \rightarrow \mathbb{B}$ ,  $P \in Prog$  and  $E = \lambda t : Tid. skip$ . The semantics*  
 348 *of a Hoare triple under partial correctness is given by:*

$$\begin{aligned}
 349 \quad & \{p\} \mathbf{Init} \{q\} = q(\Gamma_{\mathbf{Init}}) \\
 350 \quad & \{p\} P \{q\} = \forall lst, \sigma, lst', \sigma'. p(lst, \sigma) \wedge (P, lst, \sigma) \Longrightarrow^* (E, lst', \sigma') \Rightarrow q(lst', \sigma') \\
 351 \quad & \{p\} \mathbf{Init}; P \{q\} = \exists r. \{p\} \mathbf{Init} \{r\} \wedge \{r\} P \{q\}
 \end{aligned}$$

353 The classical rules of sequential Hoare logic for compound (i.e., non-atomic) statements are  
 354 given in Figure 7. Soundness of these proof rules (with respect to Definition 2) holds for  
 355 exactly the same reason as soundness of Hoare logic [15].

$$\begin{array}{c}
\text{SKIP} \frac{}{\{p\} \mathbf{skip} \{p\}} \quad \text{SEQ} \frac{\{p\} C_1 \{r\} \quad \{r\} C_2 \{q\}}{\{p\} C_1; C_2 \{q\}} \\
\\
\text{IF} \frac{\{p \wedge B\} C_1 \{q\} \quad \{p \wedge \neg B\} C_2 \{q\}}{\{p\} \mathbf{if } B \mathbf{ then } C_1 \mathbf{ else } C_2 \{q\}} \quad \text{WHILE} \frac{\{p \wedge B\} C \{p\}}{\{p\} \mathbf{while } B \mathbf{ do } C \{p \wedge \neg B\}} \\
\\
\text{UNTIL} \frac{\{p\} C \{r\} \quad \{r\} \mathbf{while } \neg B \mathbf{ do } C \{r \wedge B\}}{\{p\} \mathbf{do } C \mathbf{ until } B \{r \wedge B\}} \quad \text{CONS} \frac{p \Rightarrow p' \quad \{p'\} C \{q'\} \quad q' \Rightarrow q}{\{p\} C \{q\}}
\end{array}$$

■ **Figure 7** Classical proof rules for sequential programs

356 The sequential part is combined with the Owicki-Gries rule for concurrent composition  
 357 in the standard way [27, 7]. First, we construct *proof outlines* for every component of the  
 358 concurrent program in isolation. A proof outline inserts assertions (in  $\{ \}$  brackets) into a  
 359 program. In a so-called *standard* proof outline every statement  $R$  of the program has exactly  
 360 one assertion before it. This assertion is its *precondition*,  $\text{pre}(R)$ . Next, all assertions in one  
 361 component have to be checked for non-interference with all statements in other components.

362 ► **Definition 3.** A statement  $R \in ACom$  with precondition  $\text{pre}(R)$  (in the standard proof  
 363 outline) does not interfere with an assertion  $p$  if

$$364 \quad \{p \wedge \text{pre}(R)\} R \{p\} .$$

365 Interference freedom guarantees that proof outlines in each thread are stable under the  
 366 execution of other threads. This is formalised in the Owicki-Gries proof rule for concurrent  
 367 composition:

$$368 \quad \text{PARALLEL} \frac{\text{Proof outlines } \{p_i\} C_i \{q_i\} \text{ are interference free}}{\{\bigwedge_{i=1}^n p_i\} C_1 || \dots || C_n \{\bigwedge_{i=1}^n q_i\}}$$

370 We say a proof outline is *valid* if it is both sequentially valid (or locally correct) and  
 371 interference free.

372 Finally, there is a standard proof rule for auxiliary variables in parallel programs [7]. Let  
 373  $V$  be a set of auxiliary variables of a parallel program  $P$  and  $q$  be a predicate that does not  
 374 mention auxiliary variables. Then we can prove that a Hoare triple holds for a program  
 375 extended with auxiliary variables and transfer this proof to the original program:

$$376 \quad \text{AUXVAR} \frac{\{true\} \mathbf{Init}; P \{q\}}{\{true\} \mathbf{Init}_0; P_0 \{q\}} \text{ provided } \text{vars}(q) \cap V = \emptyset$$

377 where  $\mathbf{Init}_0$  is obtained from  $\mathbf{Init}$  by removing all auxiliary assignments and  $P_0$  is obtained  
 378 by replacing all statements  $\langle A, a := E \rangle$  in  $P$  (for  $a \in V$ ) by  $A$ .

## 379 5.2 An Assertion Language

380 We studied a number of well-known litmus tests and examples and discovered three main  
 381 categories of assertions required for specification and verification of a wide range of problems.  
 382 These three main categories are dealing with (values of) writes to variables and the order in  
 383 which they occur.

384 ■ **Observability.** Observability assertions describe if or when a thread may observe or  
 385 has encountered a write to a variable. As described in Section 2, these assertions are  
 386 thread-specific and deal with the thread's view. We repeat the main ideas here to simplify  
 387 comparison with the other types of assertions. The main observability assertions are as  
 388 follows:

- 389 1. **Possible observation** which is denoted by  $x \approx_t u$  means that thread  $t$  *may* observe  
 390 value  $u$  for  $x$ . The formal definition and an example motivating this assertion is given  
 391 in Section 5.4.
- 392 2. **Definite observation** which is denoted by  $x =_t u$  means that thread  $t$  *must* observe  
 393 the value  $u$  for  $x$ . The formal definition and an example motivating this assertion is  
 394 given in Section 5.3.
- 395 3. **Conditional observation** which is denoted by  $[x = u](y =_t v)$  means that if thread  
 396  $t$  synchronises with a write to variable  $x$  with value  $u$ , it *must* observe value  $v$  for  $y$ .  
 397 The formal definition and an example motivating this assertion is given in Section 5.4.
- 398 4. **Encountered value** which is denoted by  $x \stackrel{enc}{=} v$  means that thread  $t$  has encountered  
 399 (had the opportunity to observe) a write to variable  $x$  with value  $v$ . The formal definition  
 400 and three examples motivating this assertion are given in Section 5.5.

401 ■ **Ordering.** Ordering assertions specify the order of values written to a variable by  
 402 different writes. These assertions are thread-independent and specify an order over the  
 403 timestamp of various writes with specific values:

- 404 1. **Possible value order** which is denoted by  $m \prec_x n$  means that there exists two writes  
 405  $w$  and  $w'$  to variable  $x$  where the timestamp of  $w'$  is larger than the timestamp of  $w$   
 406 and the value of  $w$  and  $w'$  is  $m$  and  $n$ , respectively.
- 407 2. **Definite value order** which is denoted by  $m \ll_x n$  means that for all writes  $w$  and  
 408  $w'$  to  $x$  where the value of  $w$  is  $m$  and the value of  $w'$  is  $n$ , the timestamp of  $w'$  is  
 409 larger than the timestamp of  $w$  and  $m \prec_x n$ .

410 Both the above assertions are formally defined in Section 5.5 and examples showing their  
 411 usage are provided.

412 ■ **Occurrence.** Occurrence assertions specify the occurrence of a write with a specific  
 413 value to a variable (regardless of observability). Similar to the previous category, these  
 414 assertions are thread-independent:

- 415 1. **Value occurrence** assertions specify the limit of occurrence of writes to a variable  
 416 with a specific value. For instance,  $\mathbb{0}_x n$  means that no write with value  $n$  to variable  
 417  $x$  has occurred or  $\mathbb{1}_x n$  means that there is *at most* one write with value  $n$  to  $x$  in  
 418 the current state. The formal definition and examples of these assertions are given in  
 419 Section 5.5.
- 420 2. **Initial value** which is denoted by  $x_{\text{Init}} = n$  means that the initial value written to  $x$   
 421 is  $n$ . The formal definition and examples of this assertion are also given in Section 5.5.
- 422 3. **Covered write** assertions, denoted by  $\mathbf{C}_x^n$ , state that all writes to variable  $x$  except  
 423 the last write are covered by an update (see Section 4.2), and that the last write to  $x$   
 424 has value  $n$ . This assertion is formally defined in Section 6 and is used in verification  
 425 of Peterson's mutual exclusion algorithm.

### 426 5.3 Load Buffering

427 Our first example is the load buffering litmus test (see Figure 8), which we can show satisfies  
 428 the postcondition  $r1 = 0 \vee r2 = 0$  since our semantics assumes absence of cycles in the  
 429 sequence-before relation combined with reads-from [20, 12]. The assertions about the C11

**Init:**  $x := 0; y := 0; r1 := 0; r2 := 0;$   
 $\{x =_1 0 \wedge y =_2 0 \wedge r1 = 0 \wedge r2 = 0\}$   
**Thread 1**      **Thread 2**  
 $\{y =_2 0 \wedge r2 = 0\}$        $\{x =_1 0 \wedge r1 = 0\}$   
 $1 : r1 \leftarrow x;$        $3 : r2 \leftarrow y;$   
 $\{y =_2 0 \wedge r2 = 0\}$        $\{x =_1 0 \wedge r1 = 0\}$   
 $2 : y := 1;$        $4 : x := 1;$   
 $\{r1 = 0 \vee r2 = 0\}$        $\{r1 = 0 \vee r2 = 0\}$   
 $\{r1 = 0 \vee r2 = 0\}$

■ **Figure 8** Proof outline for load buffering

state capture properties about *definite observations* (i.e., observability assertions), which we formalise below.

For a set of writes  $W$  and variable  $x \in \text{Var}_G$ , let  $W_x = \{w \in W \mid \text{var}(w) = x\}$  be the set of writes in  $W$  that write to  $x$ . We define the *last write* to  $x$  in  $W$  as:

$$\text{last}(W, x) = w \iff w \in W_x \wedge (\forall w' \in W_x. \text{tst}(w') \leq \text{tst}(w))$$

Moreover, we define the definite observation of a view function, *view* with respect to a set of writes as follows:

$$\text{dview}(\text{view}, W, x) = n \iff \text{view}(x) = \text{last}(W, x) \wedge \text{wrrval}(\text{last}(W, x)) = n$$

The first conjunct ensures that the viewfront of *view* for  $x$  is the last write to  $x$  in  $W$ , and the second conjunct ensures that the value written by the last write to  $x$  in  $W$  is  $n$ .

**Definite observation.** For a variable  $x$ , thread  $t$  and value  $n$ , we define:

$$x =_t n \quad = \quad \lambda \sigma. \text{dview}(\sigma.\text{tview}_t, \sigma.\text{writes}, x) = n$$

Expanding this out, we obtain:

$$\sigma.\text{tview}_t(x) = \text{last}(\sigma.\text{writes}, x) \wedge \text{wrrval}(\text{last}(\sigma.\text{writes}, x)) = n$$

The first conjunct ensures that the viewfront of  $t$  for  $x$  is the last write to  $x$  in  $\sigma$  (thus  $t$  can only read this last write to  $x$ ). The second conjunct ensures that the value written by the last write is  $n$ . The function *dview* is also used in the definition of conditional observation in Section 5.4.

The proof of load buffering relies on the basic axioms in the following lemma. We assume *atoms*(**Init**) returns the set of assignments contained within **Init**.

► **Lemma 4.** *Each of the basic axioms below is sound (as per Definition 2), where the statements are decorated with the thread identifier of the executing thread.*

$$\begin{array}{c}
 \text{INIT} \frac{x := n \in \text{atoms}(\mathbf{Init})}{\{true\} \mathbf{Init} \{x =_t n\}} \quad \text{DOPRES-RD} \frac{}{\{x =_{t'} m\} r \xleftarrow[t]{[A]} y \{x =_{t'} m\}} \\
 \\
 \text{DOPRES-WR} \frac{x \neq y}{\{x =_{t'} n\} y :=_t m \{x =_{t'} n\}}
 \end{array}$$

**Proof.** In Isabelle. ◀

Thus by rule INIT an assignment  $x := n$  in INIT ensures that  $x =_t n$  for all threads  $t$  holds at program start. Note that such an initial assertion for the entire program is not subject to non-interference checks. The rule DOPRES-RD states that a definite observation  $x =_{t'} m$  is invariant over a read step executed by thread  $t$ . Note that pre/post conditions for DOPRES-RD refer to thread  $t'$ , while the read statement refers to thread  $t$ . Also note that there is no additional restriction on  $t$  and  $t'$ , thus the rule applies regardless of whether  $t = t'$ , or not. Similarly, there are two global variables  $x$  and  $y$  mentioned in the rule, but there are no further restrictions on their values. Rule DOPRES-WR gives a condition for invariance of a definite observation assertion over a write. It requires that the variable being observed is different from the variable that is updated.

► **Theorem 5.** *The proof outline for load buffering in Figure 8 is valid.*

**Proof.** The proof has been established in Isabelle. We outline the main steps below as it is instructive to understand the high-level proof strategy. First we establish local correctness:

- The initial condition is established by rule INIT, which is in turn used to establish the initial assertions in both threads.
- In thread 1, local correctness of the postcondition of line 1 (precondition of line 2) follows from rule DOPRES-RD, and the postcondition of line 2 follows by weakening. The proof of local correctness in thread 2 is symmetric.

We now establish interference freedom. The precondition of line 1 is interference free wrt line 3 by DOPRES-RD, and wrt line 4 by DOPRES-WR. This argument also applies to the precondition of line 2. Interference freedom of the postcondition of line 2 is trivial. The proof of interference freedom of the assertions in thread 2 is symmetric. ◀

## 5.4 Message Passing

Next we return to the message passing example from Section 2. Its verification requires the usage of the other two observability assertions.

**Possible observation.** For a variable  $x$ , thread  $t$  and value  $n$ , we define:

$$x \approx_t n \quad = \quad \lambda\sigma. \exists w \in \sigma.OW(t, x). \text{wrrval}(w) = n$$

Thus, there is a write to  $x$  that is observable to thread  $t$  with a value  $n$ .

**Conditional observation.** For variables  $x, y$ , thread  $t$  and values  $m, n$ , we define:

$$\begin{aligned} [x = n](y =_t m) \quad = \quad & \lambda\sigma. \forall w \in \sigma.OW(t, x). \text{wrrval}(w) = n \Rightarrow \\ & \text{act}(w) \in W_R \wedge \text{dview}(\sigma.\text{mview}_w, \sigma.\text{writes}, y) = m \end{aligned}$$

The antecedent assumes that the value read for  $x$  is  $n$ , and the consequent ensures that  $w$  is a releasing write such that the definite view of this write for variable  $y$  returns  $m$ . As we shall see, one useful way of establishing this condition is by falsifying the antecedent by ensuring that thread  $t$  cannot observe  $n$  for  $x$  (see (4) below).

Some useful relationships between the assertions above are given by the lemma below.

► **Lemma 6.** *For variables  $x, y \in \text{Var}_G$ , thread  $t$  and values  $m, n \in \text{Val}$ , each of the following holds:*

$$\text{wfs} \wedge x =_t n \Rightarrow x \approx_t n \tag{2}$$

$$\text{wfs} \wedge x =_t n \wedge x \approx_t m \Rightarrow n = m \tag{3}$$

$$x \not\approx_t n \Rightarrow [x = n](y =_t m) \tag{4}$$

$$x =_t n \wedge x =_{t'} m \Rightarrow n = m \tag{5}$$

497 **Proof.** In Isabelle. ◀

498 By (2), given a well-formed state any definite observation implies a possible observation,  
 499 and by (3) a definite observation must agree with a possible observation. By (4) if it is  
 500 not possible to observe the antecedent of a conditional observation, then the conditional  
 501 observation must hold. By (5) any two definite value observations must agree (since they  
 502 both observe the last write to  $x$ ).

503 The next lemma lists the basic axioms that are used to prove correctness of the message  
 504 passing example.

505 ► **Lemma 7.** *Each of the rules next is sound (as per Definition 2), where the statements are*  
 506 *decorated with the thread identifier of the executing thread.*

$$\begin{array}{l}
 507 \quad \text{MODLAST} \frac{}{\{x =_t n\} \ x :=_t m \ \{x =_t m\}} \qquad \text{MODSOME} \frac{}{\{true\} \ x :=_t m \ \{x \approx_t m\}} \\
 508 \quad \text{NPOPRES} \frac{}{\{x \not\approx_t m\} \ r \leftarrow_{t'}^{[A]} y \ \{x \not\approx_t m\}} \qquad \text{NOOW} \frac{x \neq y}{\{x \not\approx_t n\} \ y :=_{t'} m \ \{x \not\approx_t n\}} \\
 509 \quad \text{READLAST} \frac{}{\{x =_t m\} \ r \leftarrow_t x \ \{r = m\}} \\
 510 \quad \text{CO-INTRO} \frac{x \neq y}{\{y =_t m \wedge x \not\approx_{t'} n\} \ x :=_t^R n \ \{[x = n](y =_{t'} m)\}} \\
 511 \quad \text{TRANSFER} \frac{}{\{[x = n](y =_t m)\} \ r \leftarrow_t^A x \ \{r = n \Rightarrow y =_t m\}} \\
 512
 \end{array}$$

513 **Proof.** In Isabelle. ◀

514 ► **Theorem 8.** *The proof outline of message passing in Figure 3 is valid.*

515 **Proof.** The proof has been established in Isabelle. We outline the main steps below. First  
 516 we show local correctness.

- 517 ■ Using INIT we establish the precondition  $f =_1 0 \wedge f =_2 0 \wedge d =_1 0 \wedge d =_2 0$ .
- 518 ■ The precondition of the program implies the initial assertions of both threads. In thread 1,  
 519 we use (3) to establish  $f \not\approx_2 1$  since (3) is logically equivalent to

$$520 \quad wfs \wedge x =_t n \wedge n \neq m \Rightarrow x \not\approx_t m$$

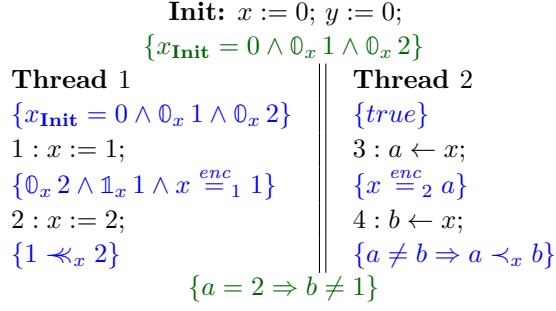
521 In thread 2, we use (3) in combination with (4).

- 522 ■ In thread 1, the post condition of line 1 (precondition of line 2) follows by application of  
 523 NOOW and MODLAST. The post condition of line 2 is trivial.
- 524 ■ In thread 2, the postcondition of line 3 follows by application of TRANSFER, while the  
 525 postcondition of line 4 follows by application of READLAST.

526 Next we show interference freedom.

- 527 ■ The preconditions of lines 1 and 2 can be shown to be interference free by applying  
 528 NPOPRES to the first conjunct and DOPRES-RD to the second.
- 529 ■ The precondition of line 3 is interference free against line 1 due to NOOW using the  
 530 existing precondition  $f \not\approx_2 1$  of line 1. The proof then follows by application of (4).  
 531 Interference freedom against line 2, is proved using CO-INTRO and the precondition at  
 532 line 2.





■ **Figure 9** Proof outline for RRC2, where  $x \in Var_G$  and  $a, b \in Var_L$

- 533 ■ The precondition of line 4 is interference free against line 1 by (5) (i.e., since the preconditions of lines 1 and 4 are contradictory). Interference freedom holds against line 2 by rule DOPRES-WR.
- 534
- 535
- 536 ■ The postconditions of lines 2 and 4 are trivially interference free.
- 537

## 5.5 Read-Read Coherence

539 Next, we verify three versions of the read-read coherence (RRC) litmus test as given in Figures 9, 10 and 11. The original RRC litmus test (Figure 10) guarantees that if one thread sees the writes to  $x$  (by threads 1 and 2) in a certain order, then the other thread see the writes in the same order. Here, the postcondition assumes that thread 3 has observed the write  $x := 1$ , then the write  $x := 2$ , while thread 4 has already seen the write  $x := 2$  when reading  $x$  at line 5. It requires that thread 4 does not subsequently see value 1 when it reads  $x$  at line 6. Figure 9 presents a simpler variation where the ordering of writes to  $x$  is enforced by the thread ordering. Figure 11 combines RRC with message passing.

547 Unlike message passing (which is a litmus test over two different variables), the RRC examples demonstrate the need for *ordering* and *occurrence* assertions which we introduce next.

550 **Possible value order.** For values  $m, n$  and variable  $x$ , we define:

$$\begin{aligned}
551 \quad m \prec_x n &= \lambda \sigma. \exists w, w' \in \sigma.writes_x. wrval(w) = m \wedge wrval(w') = n \wedge \\
552 &\quad tst(w) < tst(w')
\end{aligned}$$


553 Thus, there are two writes to  $x$  with values  $m$  and  $n$ , where the timestamp of the write with value  $m$  precedes the timestamp of the write with value  $n$ . Note that this  $m \prec_x n$  does not preclude  $n \prec_x m$ . E.g., if a thread writes  $m$  to  $x$ , then  $n$ , then  $m$  again, both  $m \prec_x n$  and  $n \prec_x m$  will hold. In this scenario,  $m \prec_x m$  also holds since there are two separate writes to  $x$  with value  $m$ .

558 **Definite value order.** For values  $m, n$  and variable  $x$ , we define:

$$\begin{aligned}
559 \quad m \ll_x n &= \lambda \sigma. (m \prec_x n)(\sigma) \wedge (\forall w, w' \in \sigma.writes_x. \\
&\quad wrval(w) = m \wedge wrval(w') = n \Rightarrow \\
560 &\quad tst(w) < tst(w'))
\end{aligned}$$

561 Note that this implies  $m \neq n$ . Unlike possible value orders if  $m \ll_x n$  holds then  $n \not\ll_x m$ . Note also that our definition allows several writes to  $x$  with values  $m$  and  $n$  provided all writes with value  $m$  occur (in timestamp order) before all writes with value  $n$ .

563



■ **Figure 10** Proof outline for RRC, where  $x \in Var_G$  and  $a, b, c, d \in Var_L$

564 **Initial value.** For values  $n$  and variable  $x$ , we define:

$$565 \quad x_{\text{Init}} = n = \lambda\sigma. \exists w \in \sigma.writes_x. wrval(w) = n \wedge$$

$$566 \quad (\forall w' \in \sigma.writes_x. w \neq w' \Rightarrow tst(w) < tst(w'))$$

567 Note that for the construction in this paper, it suffices to return the write to  $x$  with timestamp  
 568 0 since we assume that writes are initialised with timestamp 0. The definition above however,  
 569 is more robust since it also applies to situations where variables are not initialised, or  
 570 initialised to an arbitrarily chosen timestamp (as is the case in our Isabelle encoding).

571 **Encountered value.** For a variable  $x$ , thread  $t$  and value  $n$ , we define:

$$572 \quad x \stackrel{enc}{=}_t n = \lambda\sigma. \exists w \in \sigma.writes_x. tst(w) \leq tst(\sigma.tview_t(x)) \wedge wrval(x) = n$$

574 That is  $x \stackrel{enc}{=}_t n$  holds iff there is a write to  $x$  with value  $n$  whose timestamp is at most the  
 575 timestamp of the viewfront of  $t$  for  $x$ . Note that  $x \stackrel{enc}{=}_t n$  does not guarantee that  $t$  has read  
 576 the value  $n$  for  $x$ . For instance,  $x \stackrel{enc}{=}_t n$  could hold if there is a write, say  $w$ , of  $x$  with value  
 577  $n$  and  $t$  writes to  $x$  with a write whose timestamp is greater than  $tst(w)$ .

578 **Value occurrence.** These are straightforward to define in terms of our value order assertions  
 579 above. For a variable  $x$ , thread  $t$  and value  $n$ , we define:

$$580 \quad 0_x n = \exists m. x_{\text{Init}} = m \wedge m \neq n \wedge m \not\prec_x n$$

$$581 \quad 1_x n = n \not\prec_x n$$

583 Thus, if  $0_x n$  holds then there is no write with value  $n$ . If  $1_x n$  holds, then either there is no  
 584 write to  $x$  with value  $n$ , or if there is a write with value  $n$ , this is the only such write.

585 To understand the interaction between value ordering and write limit assertions, consider  
 586 the following lemma. It states that if there is a possible value order on  $x$  with  $m$  preceeding  
 587  $n$  and there is at most one write with these values, then there is a definite value order on  $x$   
 588 with  $m$  preceeding  $n$ .

589 ► **Lemma 9.** For  $x \in Var_G$  and  $m, n \in Val$ , we have:

$$590 \quad m \prec_x n \wedge 1_x m \wedge 1_x n \Rightarrow m \ll_x n \tag{6}$$

$$591 \quad m \ll_x n \Rightarrow n \not\prec_x m \tag{7}$$

593 **Proof.** In Isabelle. ◀

594 We discuss the proof of RRC2 in detail. Its proof relies on the following lemma which  
 595 captures some basic properties about value assertions.

► **Lemma 10.** *Each of the rules below is sound (as per Definition 2), where the statements are decorated with the thread identifier of the executing thread.*

$$\begin{array}{l}
\text{596} \quad \text{ZWR} \frac{m \neq n}{\{\mathbb{0}_x m\} y :=_t^{[R]} n \{\mathbb{0}_x m\}} \quad \text{DVPRES} \frac{}{\{m \prec_x n\} r \leftarrow_t^{[A]} y \{m \prec_x n\}} \\
\text{599} \quad \text{1INTRO} \frac{i \neq m}{\{x_{\text{Init}} = i \wedge \mathbb{0}_x m\} x :=_t^{[R]} m \{\mathbb{1}_x m\}} \quad \text{ENCWR} \frac{}{\{true\} x :=_t^{[R]} m \{x \stackrel{enc}{=} m\}} \\
\text{600} \quad \text{ENCRD} \frac{}{\{true\} r \leftarrow_t^{[A]} x \{x \stackrel{enc}{=} r\}} \quad \text{EPO} \frac{}{\{x \stackrel{enc}{=} m\} r \leftarrow_t^{[A]} x \{r \neq m \Rightarrow m \prec_x r\}} \\
\text{601} \quad \text{DVINTRO} \frac{i \neq n}{\{x_{\text{Init}} = i \wedge \mathbb{0}_x n \wedge \mathbb{1}_x m \wedge x \stackrel{enc}{=} m\} x :=_t^{[R]} n \{m \prec_x n\}} \\
\text{602} \quad \text{1PRESR} \frac{}{\{\mathbb{1}_x m\} r \leftarrow_t^{[A]} y \{\mathbb{1}_x m\}} \quad \text{POR} \frac{}{\{m \prec_x n\} C \{m \prec_x n\}} \\
\text{603}
\end{array}$$

► **Proof.** In Isabelle. ◀

► **Theorem 11.** *The proof outline for RRC2 in Figure 9 is valid.*

► **Proof.** This proof has been mechanised in Isabelle. Once again, we describe the proof outline to give an overview of how our proofs are used. For local correctness we have the following.

- The initialisation clearly satisfies the precondition of the program, and this implies the precondition of thread 1. The precondition of thread 2 is trivial.
- Next we consider the postcondition of line 1. The first conjunct holds by ZWR, the second conjunct holds by 1INTRO and the third by rule ENCWR.
- The postcondition of line 2 holds by rule DVINTRO.
- In thread 2, the postcondition of line 3 holds by rule ENCRD, and the postcondition of line 4 holds by rule EPO.

Next we check interference freedom.

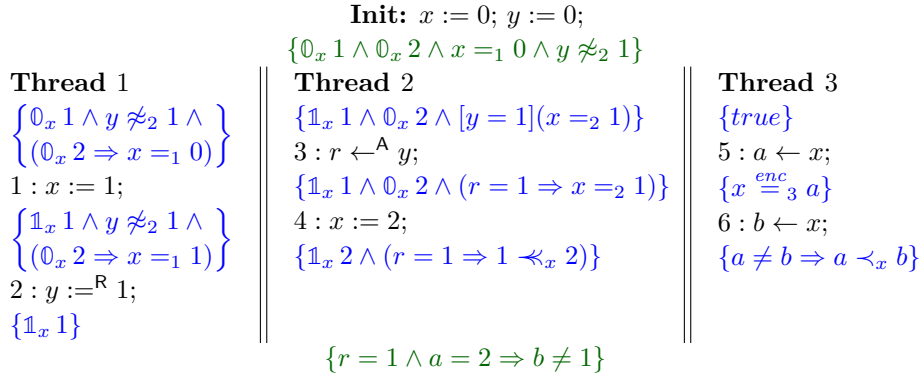
- The precondition of line 1 is stable with respect to lines 3 and 4 by ZWR.
- Next consider the precondition of line 2. The first and second conjuncts are stable with respect to lines 3 and 4 by ZWR and 1PRESR, respectively. The third conjunct is trivially preserved (see Isabelle).
- The postcondition of line 2 holds by DVPRES.
- The precondition of line 3 is trivial and the postcondition of line 3 holds by POR.

Correctness of RRC and RRC3 is established by the following theorem.

► **Theorem 12.** *The proof outlines for RRC and RRC3 in Figure 10 and Figure 11, respectively are valid.*

► **Proof.** In Isabelle. ◀

For RRC (Figure 10), the precondition of line 4 records the fact that thread 3 has encountered  $a$  (whatever the value of  $a$  may be). Moreover, it guarantees that there is at most one write of  $x$  with values 1 and 2. The first conjunct (i.e.,  $x \stackrel{enc}{=} a$ ) allows us to conclude that after  $x$  is read at line 4, if  $a$  and  $b$  are different, then the value for  $a$  is



■ **Figure 11** Proof outline for RRC3, where  $x, y \in Var_G$  and  $a, b \in Var_L$

possibly ordered before the value for  $b$ . The second and third conditions are used to establish the postconditions  $1_x 1$  and  $1_x 2$ . This argument also applies to the assertions in thread 4. Finally, we show that the postcondition of the program holds as follows, where we assume  $post$  is the conjunction of the postcondition of each thread.

$$\begin{aligned}
& post \Rightarrow (a = 1 \wedge b = 2 \wedge c = 2 \Rightarrow d \neq 1) \\
\iff & post \wedge a = 1 \wedge b = 2 \wedge c = 2 \wedge d = 1 \Rightarrow false & \text{(logic)} \\
\iff & 1_x 1 \wedge 1_x 2 \wedge 1 \prec_x 2 \wedge 2 \prec_x 1 \Rightarrow false & \text{(logic)} \\
\iff & 1 \prec_x 2 \wedge 2 \prec_x 1 \Rightarrow false & (6) \\
\iff & true & (7)
\end{aligned}$$

The calculation above has been verified with Isabelle, but we recall the proof here as it provides insight into the interactions between different value assertions.

RRC3 (Figure 11) combines message passing on  $y$  with RRC on  $x$ . Namely, knowledge of  $x := 1$  in thread 1 is transferred to thread 2 using a release-acquire synchronisation on  $y$ . Thus, if thread 2 reads 1 for  $y$  it must also have encountered 1 for  $x$ . Thus, if  $r = 1$ , then the write on line 4 must have happened *after* the write on line 1. This means that it should be impossible for thread 3 to read 2 for  $x$  (at line 5) then read 1 for  $x$  (at line 6). Unlike message passing, in RRC3, the “data” variable  $x$  is updated both before and after synchronisation. Thus, the assertions on definite values (e.g.,  $x =_1 1$ ) become conditional on whether line 4 has already been executed. In particular, the antecedent  $0_x 2$  allows us to assume that line 4 has not yet been executed. As with RRC, we must separately prove that the conjunction of the postconditions of the threads implies the postcondition of the program. This proof is mechanised in Isabelle, and is elided here.

## 6 Case study: Peterson’s algorithm

We turn to our final case study, the verification of the mutual exclusion property of a version of Peterson’s algorithm. The complexity of this case study is much greater than our earlier examples. This program contains a loop, features a careful mixture of relaxed and release/acquire operations to the same variable, and an RMW operation whose precise semantics is critical to the correctness of the algorithm.

**Init:**  $flag_1 := 0; flag_2 := 0; turn := 0 \wedge after_1 := false; after_2 := false$

**Thread 1**

$$\left\{ \neg after_1 \wedge flag_1 =_1 0 \wedge turn \not\approx_2 2 \wedge (C_{turn}^0 \vee [turn = 1](flag_2 =_1 1)) \right\}$$

$$\left\{ \wedge(after_2 \Rightarrow C_{turn}^1 \wedge [turn = 1](flag_2 =_1 1)) \right\}$$

1:  $flag_1 := 1$  ;

$$\left\{ \neg after_1 \wedge flag_1 =_1 1 \wedge turn \not\approx_2 2 \wedge (after_2 \Rightarrow C_{turn}^1 \wedge [turn = 1](flag_2 =_1 1)) \right\}$$

2:  $\langle turn.\mathbf{swap}(2)^{RA} ; after_1 := true \rangle$

$$\left\{ after_1 \wedge (after_2 \wedge (flag_2 \approx_1 0 \vee turn \approx_1 1) \Rightarrow turn =_2 1) \right\}$$

**do**

3:  $r_1 \leftarrow^A flag_2$

$$\left\{ after_1 \wedge (after_2 \wedge (r_1 = 0 \vee turn \approx_1 1 \vee flag_2 \approx_1 0) \Rightarrow turn =_2 1) \right\}$$

4:  $r_2 \leftarrow turn$

$$\left\{ after_1 \wedge (after_2 \wedge (r_1 = 0 \vee r_2 = 1 \vee turn \approx_1 1 \vee flag_2 \approx_1 0) \Rightarrow turn =_2 1) \right\}$$

5: **until**  $(r_1 = 0 \vee r_2 = 1)$

$$\left\{ after_1 \wedge (after_2 \Rightarrow turn =_2 1) \right\}$$

6: Critical section ;

7:  $\langle flag_1 :=^R 0 ; after_1 := false \rangle$

■ **Figure 12** Peterson’s algorithm (adapted from [36]) and its proof outline. **Thread 2** (not shown) is symmetric.

Our version of Peterson’s algorithm<sup>4</sup>, presented in Figure 12 is a mutual exclusion algorithm for two threads implemented for C11 using release-acquire annotations [36]. As with the original algorithm, variable  $flag_i$ , for  $i \in \{1, 2\}$  is used to indicate whether thread  $i$  intends to enter its critical section. In this version of the algorithm, we let  $flag_i$  range over  $\{0, 1\}$ , where 0 is used for the boolean value “false”, and 1 is used for the boolean value “true”. The shared variable  $turn$  is used to cause a thread to “give way” when both threads intend to enter their critical sections at the same time. Our verification uses auxiliary variables  $after_i$  for each thread  $i$  (as does the proof for a sequentially consistent setting in [7]), the purpose of which we describe below.

We describe the algorithm for thread 1; the other thread is symmetric. For now, we ignore the assertions. The flag variable is set to 1 (line 1) using a relaxed write (which cannot induce any synchronisation), but is set to 0 (line 7) using a release annotation. The intention of the latter is to synchronise this write (of 0 to  $flag_1$ ) with the read of  $flag_1$  at line 3 in thread 2. The value of  $turn$  is set using a **swap** command. The **swap** is implemented using an C11 RMW operation that has both the release and acquire annotations. When the **swap** is executed, as part of the same transition, the auxiliary variable  $after_1$  is also set, indicating that thread 1 is ready to enter the busy wait loop beginning at line 3, and then to enter the critical section.

The busy wait loop forces thread 0 to wait until either  $flag_2$  is 0 (indicating that thread 2 is not trying to enter the critical section) or  $turn = 1$  (indicating that it is thread 1’s turn to enter the critical section). Note that the read of  $turn$  within the guard of the busy wait loop (line 5) is relaxed.

We turn now to the proof that this version of Peterson’s algorithm has the mutual exclusion property. We prove mutual exclusion in two steps. First, we show that the given proof outline is valid, and second, that the conjunction of the precondition of thread 1’s critical section (line 6) and thread 2’s must be false. Therefore, the two threads cannot

<sup>4</sup> For simplicity our version of the algorithm does not have an outermost loop.

686 simultaneously be in their critical sections.

687 We deal with the second step first by showing that the formula below is *false*:

$$688 \text{ } after_1 \wedge (after_2 \Rightarrow turn =_2 1) \wedge after_2 \wedge (after_1 \Rightarrow turn =_1 2)$$

690 It is easy to see that this implies  $turn =_1 2 \wedge turn =_2 1$ . However, by (5) this situation is  
691 impossible.

692 The first step is more elaborate and we only describe certain aspects. The precondition of  
693 line 3 is also an invariant of the busy wait loop. This assertion ensures that if thread 1 is able  
694 to exit the busy wait loop, then the precondition of the critical section will be satisfied. Note  
695 that thread 1 exits the loop if it reads 0 from  $flag_2$  (which is only possible when  $flag_2 \approx_1 0$ )  
696 or it reads 1 from  $turn$  (which is only possible when  $turn \approx_1 1$ ). The invariant states that if  
697 one of these conditions holds in a state where thread 2 is waiting to enter the critical section  
698 (that is,  $after_2$ ), we can conclude  $turn =_2 1$  as required.

699 Proving that the precondition of line 3 is satisfied in the post-state of line 2 requires using  
700 a feature of our assertion language, closely related to the semantics of RMW operations,  
701 that we now introduce. Recall from the UPDATE rule in Figure 5 that whenever a write  $w$   
702 is read-from by an RMW operation,  $w$  becomes *covered*, so that no later write (or RMW)  
703 operation can be inserted between  $w$  and the RMW. This feature of C11 is critical to the  
704 correctness of Peterson's algorithm. Observe that the  $turn$  variable is only modified by RMW  
705 operations, and therefore every write to  $turn$  is covered, except the last. To formally state  
706 this, we need the third *occurrence* assertion  $\mathbf{C}_x^n$ , defined as follows.

$$707 \quad \mathbf{C}_x^n = \lambda \sigma. \forall w \in \sigma.writes_x. w \notin \sigma.covered \Rightarrow wrval(w) = n \wedge w = last(W, x)$$

709 So  $\mathbf{C}_x^n$  means that every write to  $x$  except the last is covered and the value written by that  
710 last write is  $n$ .

711 We use the following lemma on covered.

► **Lemma 13.**

$$712 \quad \text{CVD-UPD} \frac{}{\{\mathbf{C}_x^n\} \text{ upd}^{\text{RA}}(x, m, l) \{m = n \wedge \mathbf{C}_x^l\}} \quad \text{CVD-WR} \frac{x \neq y}{\{\mathbf{C}_x^n\} y :=^{[R]} m \{\mathbf{C}_x^n\}}$$

$$713 \quad \text{CVD-RD} \frac{}{\{\mathbf{C}_x^n\} r \leftarrow^{[A]} y \{\mathbf{C}_x^n\}} \quad \text{CVD-DOBS} \frac{}{\{\mathbf{C}_x^n\} \text{ upd}^{\text{RA}}(x, m, n) \{x =_t n\}}$$

715 Rule CVD-UPD states that if  $\mathbf{C}_x^n$  holds in the pre-state, then after executing  $\text{upd}^{\text{RA}}(x, m, l)$ ,  
716 we must have that  $m = n$  (since the only value available for the RMW to read is  $n$ ), and  
717 furthermore we obtain a new covered predicate  $\mathbf{C}_x^l$ . Thus, it is possible to maintain a covered  
718 predicate in a program (with possibly different return values) by ensuring each modification to  
719 the covered variable is via an RMW. This is a property that is true of Peterson's algorithm  
720 as given in Figure 12. Rules CVD-WR and CVD-RD give preservation properties for the  
721 covered assertion for a read and a write, respectively. Finally, CVD-DOBS is used to establish  
722 a definite observation of a covered assertion after an update command.

723 The precondition of line 2 asserts that if thread 2 is ready to enter the critical section  
724 (that is,  $after_2$ ) then the RMW to be executed at line 2 must read from the last write which  
725 has value 1 (that is,  $\mathbf{C}_{turn}^1$ ) and when this RMW occurs then thread 1 will definitely see  
726  $flag_2$  set (that is,  $[turn = 1](flag_2 =_1 1)$ ). This is enough to show that if  $after_2$  then in the  
727 post-state of the RMW,  $flag_2 \not\approx_1 0$  which is sufficient to prove the postcondition of line 2.

728 Of course, the sequential reasoning above must be combined with an interference freedom  
729 check, which is supported by a set of basic lemmas describing how  $\mathbf{C}_x^n$  is updated. This leads  
730 to the following theorem, which establishes validity of the proof outline.

```

lemma d_obs_Wr_set:
  assumes "wfs  $\sigma$ "
    and "wr_val a = Some n"
    and "avar a = x"
    and " $[x =_t m] \sigma$ "
    and "step t a  $\sigma \sigma'$ "
  shows " $[x =_t n] \sigma'$ "

corollary d_obs_WrX_set:
  " $wfs \sigma \implies [x =_t m] \sigma \implies \sigma [x := n]_t \sigma' \implies [x =_t n] \sigma'$ "

corollary d_obs_WrR_set :
  " $wfs \sigma \implies [x =_t m] \sigma \implies \sigma [x :=^R n]_t \sigma' \implies [x =_t n] \sigma'$ "

corollary d_obs_RMW_set :
  " $wfs \sigma \implies [x =_t m] \sigma \implies \sigma \text{RMW}[x,w,n]_t \sigma' \implies [x =_t n] \sigma'$ "

```

■ **Figure 13** Isabelle encoding of basic axioms over C11 assertions

731 ► **Theorem 14.** *The proof outline of Peterson’s algorithm (Figure 12) is valid.*

732 **Proof.** In Isabelle. ◀

733 We note that Peterson’s algorithm represents a challenge in deductive verification. Unlike  
 734 the litmus tests presented above, there is sufficient complexity in the algorithm and the  
 735 resulting proof outline so that pen-and-paper proofs cannot be trusted. Using our mech-  
 736 anisation, we explored several variations of the proof outline in Figure 12, and discovered  
 737 simplifications to our original pen-and-paper proofs.

## 738 7 Mechanisation

739 As already mentioned, the operational semantics as well as all lemmas and theorems presented  
 740 in this paper have been mechanised in Isabelle. In this section, we discuss our mechanisation  
 741 effort.

742 To prove the lemmas about basic assertions, we typically prove a more general result  
 743 relating to reads and writes, which are then specialised so that they can be used in the  
 744 verification of the algorithms. For example, we first prove the lemma in Figure 13, which  
 745 describes changes to definite values and applies to any writing transition. This is then  
 746 specialised to the corollaries on the right, which are easier for Isabelle to find when performing  
 747 the verification of the proof outlines.

748 The generic lemmas require some amount of interactive work. However, once verified, it is  
 749 straightforward to use them to prove the corollaries. For example, corollary `d_obs_WrX_set` in  
 750 Figure 13 is verified with the command “`by (metis WrX_def avar.simps(2) d_obs_Wr_set`  
 751 `wr_val.simps(1))`”, which is found automatically by Isabelle’s built in `sledgehammer`  
 752 tool [10].

753 Such lemmas and corollaries are in turn used in the proofs of programs. First the program  
 754 state (i.e.,  $\Sigma_{C11}$ ) is encoded as a `record` type with a special variable that models the C11  
 755 state. The programs themselves are encoded as a relation over these records with program  
 756 counters modelling control flow. This allows the proof outlines to be encoded as predicates



mapping program counters to the assertions at that control point. We then verify a set of lemmas that guarantee local correctness and interference freedom, where we decompose proofs and apply case analysis over the individual program steps (e.g., reads, writes for each thread). Once a proof has been decomposed, `sledgehammer` is able to find the relevant corollaries (e.g., those in Figure 13) to discharge proofs automatically.

## 8 Related Work

The semantics and verification of programs running on weak memory models has recently received a lot of attention. Lahav [21] gives a brief survey for C11.

Our timestamp based operational semantics is motivated by ideas in [13] and is similar to the semantics of Kaiser et al. [16, 17]. We note there are differences in coverage of the memory models in [13, 16, 17]. Dolan et al. [13] cover a sequentially consistent (SC) and relaxed accesses for OCAML, where the SC operations behave like Java volatiles. Kaiser et al [16] covers non-atomics and release-acquire, while Kang et al. [17] support a much larger fragment of C11, including so-called load-buffering cycles.

Abdulla et al. have shown the reachability problem for release-acquire to be *undecidable* [2]. A number of works target *model checking* for weak memory, e.g., by explicitly encoding architectural structures leading to weak behaviour, like store buffers [33, 5]. Ponce de León et al. [30, 14] have developed a bounded model checker for weak memory models, taking the axiomatic description of a memory model as input. (Bounded) model checkers for specific weak memory models are furthermore the tools CBMC [6] (for TSO), NIDHUGG [1] (for TSO and PSO), RCMC [18] (for C11) and GENMC [19] (again, parametric in memory model).

A (non-automatic) reasoning technique for proving invariants – parameterised by a weak memory model – has been proposed by Alglave and Cousot [4]. They propose a new semantics, different from an operational one without any coherence order (or modification order) constraining the order of writes to memory. Their assertions contain so-called pythia variables to uniquely identify values of read events, and require a separate communication proof (differentiating their method from standard Owicki-Gries reasoning). They say “In addition to the initialisation, sequential, and non-interference proof, the main difference with Owicki and Gries [27] (and Lamport 1977) is the use of pythia variables and the read-from relation in assertions and the communication proof showing that reads-from is well-formed.” [4]. Our method in contrast only requires the initialisation, sequential, and non-interference proofs as with the original technique.

Another manual method for the RC11 memory model has been developed by Doherty et al. [12], who cover the message passing example and Peterson’s algorithm. Our work is inspired by this existing work, however, there are several differences. They use a classical model of the C11 state (expressed in terms of a set of relations, e.g., reads-from, sequenced-before etc), develop assertions over these relations and a small proof calculus for these assertions. However, their methods are at a lower level of abstraction than the techniques presented in this paper since the assertions are stated in terms of individual relations that make up each state. Thus, it is not possible to directly develop a Hoare logic for their assertions and mechanisation itself is more difficult.

Also close to our work is that of Lahav and Vafeiadis [22] who also develop an Owicki-Gries style proof calculus. We consider all their examples except RCU — our logic can handle the RCU example, but this proof has thus far not been mechanised. Moreover, we include several other case studies such as litmus tests that combine read-read coherence with message passing and the non-trivial Peterson’s algorithm. There are several additional differences to note. (1)

803 Lahav and Vafeiadis’ proof calculus is developed in the absence of an operational semantics,  
 804 and hence, their definition of a valid Hoare triple is non standard (see [22, Definition 9]). A  
 805 consequence of this is that they must be careful about the introduction of auxiliary variables,  
 806 resorting to the more restricted notion of a *ghost* variable. In contrast, we use traditional  
 807 auxiliary variables — an auxiliary variable must not affect the control flow of a program  
 808 nor be assigned to any program variable. Note however, that to simplify the presentation, we  
 809 use auxiliary variables in a more restricted manner (see Section 3). (2) They do not handle  
 810 relaxed accesses — as stated in their conclusion: “While OGRA’s non-interference condition  
 811 appears to be restrictive, it is unsound for weaker memory models, such as C11’s relaxed  
 812 accesses ...”. (3) They do not provide a mechanisation.

813 A frequently employed starting point for program logic is separation logic, for which  
 814 a number of extensions to weak memory exist (GPS [34], RSL [16]). Svendsen et al. [32]  
 815 propose a separation logic based on the promising semantics of Kang et al. [17]. The principle  
 816 of ownership transfer used therein naturally fits to message passing using release acquire.  
 817 Prover support for such separation logic based proofs — like ours with Isabelle — has been  
 818 developed for the Iris proof system [16]. Tool support has also been developed by Summers  
 819 and Müller [31], where the RSL logic has been encoded in the Viper tool, offering a level  
 820 of proof automation. Their encoding is proved sound and complete with respect to RSL.  
 821 However, such efforts do not provide a clear link between C11 semantics and traditional  
 822 reasoning using Hoare logics.

## 823 9 Conclusion

824 In this paper, we have introduced an assertion language for C11 RAR which allows to re-use  
 825 the entire Owicki-Gries proof calculus except for the axiom of assignment. The assertion  
 826 language is based on an operational semantics for C11 RAR which we have shown to be sound  
 827 wrt. standard axiomatic semantics. We have exemplified reasoning on a number of standard  
 828 C11 RAR litmus tests as well as a C11 RAR annotated version of Peterson’s algorithm. All  
 829 proofs ranging are mechanised within Isabelle — this includes soundness of the basic axioms  
 830 for weak memory reads, writes and updates, and validity of proof outlines for the examples  
 831 presented.

832 We are currently integrating this work<sup>5</sup> into the standard Owicki-Gries library that  
 833 is included in the Isabelle distribution [26]. As future work, we aim to tackle fragments  
 834 of C11 larger than C11 RAR, e.g., fragments that allow the load buffering example to  
 835 terminate with postcondition  $r1 = 1 \wedge r2 = 1$  [9, 17], SC annotations [20], as well as  
 836 release sequences and fences [8]. Extending our operational semantics to handle the final  
 837 two features is straightforward, but is not considered in this paper as it complicates the  
 838 semantics and detracts from our main contribution, i.e., a simple extension to Hoare logic to  
 839 enable reasoning about C11 programs. Hoare-style reasoning that incorporates the other two  
 840 features is currently being investigated.

## 841 — References —

- 842 1 Parosh Aziz Abdulla, Stavros Aronis, Mohamed Faouzi Atig, Bengt Jonsson, Carl Leonardsson,  
 843 and Konstantinos Sagonas. Stateless model checking for TSO and PSO. *Acta Inf.*, 54(8):789–  
 844 818, 2017.

<sup>5</sup> A set of preliminary results is available <https://www.dropbox.com/sh/4yr2w7792qwyw09/AACsWUXtZbK3PvyfJkqqjyDYa> within the file OG-Isabelle.zip.

- 2 Parosh Aziz Abdulla, Jatin Arora, Mohamed Faouzi Atig, and Shankara Narayanan Krishna. Verification of programs under the release-acquire semantics. In McKinley and Fisher [25], pages 1117–1132.
- 3 J. Alglave, L. Maranget, and M. Tautschnig. Herding cats: Modelling, simulation, testing, and data mining for weak memory. *ACM Trans. Program. Lang. Syst.*, 36(2):7:1–7:74, 2014.
- 4 Jade Alglave and Patrick Cousot. Ogre and Pythia: an invariance proof method for weak consistency models. In Castagna and Gordon [11], pages 3–18.
- 5 Jade Alglave, Daniel Kroening, Vincent Nimal, and Michael Tautschnig. Software verification for weak memory via program transformation. In M. Felleisen and P. Gardner, editors, *ESOP*, volume 7792 of *LNCS*, pages 512–532. Springer, 2013.
- 6 Jade Alglave, Daniel Kroening, and Michael Tautschnig. Partial orders for efficient bounded model checking of concurrent software. In Natasha Sharygina and Helmut Veith, editors, *CAV*, volume 8044 of *LNCS*, pages 141–157. Springer, 2013.
- 7 Krzysztof R. Apt, Frank S. de Boer, and Ernst-Rüdiger Olderog. *Verification of Sequential and Concurrent Programs*. Texts in Computer Science. Springer, 2009.
- 8 M. Batty, S. Owens, S. Sarkar, P. Sewell, and T. Weber. Mathematizing C++ concurrency. In T. Ball and M. Sagiv, editors, *POPL*, pages 55–66. ACM, 2011.
- 9 Mark Batty, Alastair F. Donaldson, and John Wickerson. Overhauling SC atomics in C11 and OpenCL. In *POPL*, pages 634–648. ACM, 2016.
- 10 Sascha Böhme and Tobias Nipkow. Sledgehammer: Judgement day. In *IJCAR*, volume 6173 of *Lecture Notes in Computer Science*, pages 107–121. Springer, 2010.
- 11 Giuseppe Castagna and Andrew D. Gordon, editors. *POPL*. ACM, 2017.
- 12 Simon Doherty, Brijesh Dongol, Heike Wehrheim, and John Derrick. Verifying C11 programs operationally. In Jeffrey K. Hollingsworth and Idit Keidar, editors, *PPoPP*, pages 355–365. ACM, 2019.
- 13 Stephen Dolan, KC Sivaramakrishnan, and Anil Madhavapeddy. Bounding data races in space and time. In *PLDI*, PLDI 2018, pages 242–255, New York, NY, USA, 2018. ACM.
- 14 Natalia Gavrilenko, Hernán Ponce de León, Florian Furbach, Keijo Heljanko, and Roland Meyer. BMC for weak memory models: Relation analysis for compact SMT encodings. In Isil Dillig and Serdar Tasiran, editors, *CAV*, volume 11561 of *LNCS*, pages 355–365. Springer, 2019. doi:10.1007/978-3-030-25540-4\_19.
- 15 C. A. R. Hoare. An axiomatic basis for computer programming. *Commun. ACM*, 12(10):576–580, 1969.
- 16 Jan-Oliver Kaiser, Hoang-Hai Dang, Derek Dreyer, Ori Lahav, and Viktor Vafeiadis. Strong logic for weak memory: Reasoning about release-acquire consistency in Iris. In Peter Müller, editor, *ECOOP*, volume 74 of *LIPIcs*, pages 17:1–17:29. Schloss Dagstuhl - Leibniz-Zentrum fuer Informatik, 2017.
- 17 Jeehoon Kang, Chung-Kil Hur, Ori Lahav, Viktor Vafeiadis, and Derek Dreyer. A promising semantics for relaxed-memory concurrency. In Castagna and Gordon [11], pages 175–189.
- 18 Michalis Kokologiannakis, Ori Lahav, Konstantinos Sagonas, and Viktor Vafeiadis. Effective stateless model checking for C/C++ concurrency. *PACMPL*, 2(POPL):17:1–17:32, 2018.
- 19 Michalis Kokologiannakis, Azalea Raad, and Viktor Vafeiadis. Model checking for weakly consistent libraries. In McKinley and Fisher [25], pages 96–110.
- 20 O. Lahav, V. Vafeiadis, J. Kang, C.-K. Hur, and D. Dreyer. Repairing sequential consistency in C/C++11. In *PLDI*, pages 618–632. ACM, 2017.
- 21 Ori Lahav. Verification under causally consistent shared memory. *SIGLOG News*, 6(2):43–56, 2019.
- 22 Ori Lahav and Viktor Vafeiadis. Owicki-Gries reasoning for weak memory models. In Magnús M. Halldórsson, Kazuo Iwama, Naoki Kobayashi, and Bettina Speckmann, editors, *ICALP*, volume 9135 of *LNCS*, pages 311–323. Springer, 2015.
- 23 L. Lamport. How to make a multiprocessor computer that correctly executes multiprocess programs. *IEEE Trans. Computers*, 28(9):690–691, 1979.

- 897   24   Nancy A Lynch. *Distributed algorithms*. Elsevier, 1996.
- 898   25   Kathryn S. McKinley and Kathleen Fisher, editors. *PLDI*. ACM, 2019.
- 899   26   Tobias Nipkow and Leonor Prensa Nieto. Owicki/Gries in Isabelle/HOL. In *FASE*, volume  
900   1577 of *Lecture Notes in Computer Science*, pages 188–203. Springer, 1999.
- 901   27   Susan S. Owicki and David Gries. An axiomatic proof technique for parallel programs I. *Acta*  
902   *Inf.*, 6:319–340, 1976.
- 903   28   Lawrence C. Paulson. *Isabelle - A Generic Theorem Prover (with a contribution by T. Nipkow)*,  
904   volume 828 of *LNCS*. Springer, 1994.
- 905   29   Anton Podkopaev, Ilya Sergey, and Aleksandar Nanevski. Operational aspects of C/C++  
906   concurrency. *CoRR*, abs/1606.01400, 2016. [arXiv:1606.01400](https://arxiv.org/abs/1606.01400).
- 907   30   Hernán Ponce de León, Florian Furbach, Keijo Heljanko, and Roland Meyer. BMC with  
908   memory models as modules. In Nikolaj Bjørner and Arie Gurfinkel, editors, *FMCAD*, pages  
909   1–9. IEEE, 2018.
- 910   31   Alexander J. Summers and Peter Müller. Automating deductive verification for weak-memory  
911   programs. In Dirk Beyer and Marieke Huisman, editors, *TACAS*, volume 10805 of *LNCS*,  
912   pages 190–209. Springer, 2018.
- 913   32   Kasper Svendsen, Jean Pichon-Pharabod, Marko Doko, Ori Lahav, and Viktor Vafeiadis. A  
914   separation logic for a promising semantics. In Amal Ahmed, editor, *ESOP*, volume 10801 of  
915   *LNCS*, pages 357–384. Springer, 2018.
- 916   33   Oleg Travkin, Annika Mütze, and Heike Wehrheim. SPIN as a linearizability checker under  
917   weak memory models. In Valeria Bertacco and Axel Legay, editors, *HVC*, volume 8244 of  
918   *LNCS*, pages 311–326. Springer, 2013.
- 919   34   Aaron Turon, Viktor Vafeiadis, and Derek Dreyer. GPS: navigating weak memory with ghosts,  
920   protocols, and separation. In Andrew P. Black and Todd D. Millstein, editors, *OOPSLA*,  
921   pages 691–707. ACM, 2014.
- 922   35   John Wickerson, Mark Batty, Tyler Sorensen, and George A. Constantinides. Automatically  
923   comparing memory consistency models. In Castagna and Gordon [11], pages 190–204. URL:  
924   <http://dl.acm.org/citation.cfm?id=3009838>.
- 925   36   A. Williams. [https://www.justsoftwaresolutions.co.uk/threading/petersons\\_lock\\_](https://www.justsoftwaresolutions.co.uk/threading/petersons_lock_with_C++0x_atomics.html)  
926   [with\\_C++0x\\_atomics.html](https://www.justsoftwaresolutions.co.uk/threading/petersons_lock_with_C++0x_atomics.html), 2018. Accessed: 2018-06-20.

## A Correctness of the Operational Semantics

In this section, we prove that the operational semantics presented in this paper is sound w.r.t. a standard version of the C11 RAR axiomatic semantics.

We base this proof on the operational semantics for C11 RAR developed by Doherty et al. [12] which they have proved to be equivalent to a suitable fragment of the C11 RAR axiomatic semantics of [9]. In what follows, we refer to the semantics of [12] as the *Explicit semantics* as the states of its memory semantics are simply consistent C11 RAR executions. We refer to the semantics presented in this paper as the *View semantics*.

In Section A.1, we describe the operational semantics of [12]. In Section A.2, we present a simulation relation from the View memory semantics to the Explicit memory semantics. This is sufficient to show that all sequences of operations accepted by the View semantics are also accepted by the Explicit semantics.

In what follows, we refer to states of the Explicit semantics using the variables  $E, E'$  and states of the View semantics as  $V, V'$ .

### A.1 The Explicit Memory Semantics

We refer the reader to [12] for a full discussion of the Explicit semantics. In this paper, we present only the semantics that relates to memory operations, and we do so only briefly.

States of the Explicit semantics are simply C11 RAR executions, of the form  $E = (X, \text{sb}, \text{rf}, \text{mo})$ . Where  $\text{sb}, \text{rf}, \text{mo}$  are relations on the set of operations  $X$  with their usual meanings. That is  $\text{sb} \subseteq X \times X$  is the *sequenced-before* relation;  $\text{rf} \subseteq W \times R \cap X \times X$  is the *reads-from* relation; and  $\text{mo} \subseteq W \times W \cap X \times X$  is the *modification order*.

In an Explicit state  $E = (X, \text{sb}, \text{rf}, \text{mo})$  we let  $X \subseteq \text{Act} \times \mathbb{Q} \times T$ , where  $\text{Act}$  is the set of operations and  $T$  is the set of threads. In the original presentation, [12] we specify  $X \subseteq \text{Act} \times G \times T$  where  $G$  is a set of tags that are not further specified, which serves to distinguish repeated occurrences of the same operation. Here, we let  $G = \mathbb{Q}$ , for uniformity with the view semantics.

The Explicit semantics uses additional *synchronized-with* (denoted  $\text{sw}$ ) and *happens-before* (denoted  $\text{hb}$ ) relations, defined as follows:

$$\text{sw} = \text{rf} \cap (W_R \times R_A) \quad (8)$$

$$\text{hb} = (\text{sb} \cup \text{sw})^+ \quad (9)$$

In the Explicit semantics, all variables are initialised by a special *initialising thread*  $0 \in T$ . Define the set of *initialising writes* to be  $\text{IW}r = \{w \in W \mid \text{tid}(w) = 0\}$ . The initial states of our operational model are those of the form  $E_0 = ((I, \emptyset), \emptyset, \emptyset)$  where  $I \subseteq \text{IW}r$ , and for each variable  $x$ , there is exactly one write  $w \in I$  such that  $\text{var}(w) = x$ . For a state  $E = ((X, \_), \_, \_)$ , let  $I_E = X \cap \text{IW}r$ .

The relation  $\text{fr} = (\text{rf}^{-1}; \text{mo}) \setminus \text{Id}$  (where  $;$  is relational composition) is the “from-read” relation, that relates each read to all writes that are  $\text{mo}$ -after the write the read has read from. We must subtract  $\text{Id}$  (identity) edges from  $\text{rf}^{-1}; \text{mo}$  to cope with update events, which have the potential to induce reflexivity in  $\text{fr}$  [9, 20].

In addition, the semantics uses the *extended coherence order* [20], denoted  $\text{eco}$ , which fixes the order of reads and writes to each variable. Formally we define:

$$\text{eco} = (\text{fr} \cup \text{mo} \cup \text{rf})^+ \quad (10)$$

$$\begin{array}{c}
\text{READ} \frac{a \in \{rd(x, n), rd^A(x, n)\} \quad wrval(w) = n}{w \in E.OW(t, x) \quad rf' = rf \cup \{(w, e)\} \quad mo' = mo} \\
\hline
((X, sb), rf, mo) \xrightarrow{e} ((X, sb) + e, rf', mo') \\
\\
\text{WRITE} \frac{a \in \{wr(x, n), wr^R(x, n)\} \quad w \in E.OW(t, x) \setminus E.covered}{rf' = rf \quad mo' = mo[w, e]} \\
\hline
((X, sb), rf, mo) \xrightarrow{e} ((X, sb) + e, rf', mo') \\
\\
\text{RMW} \frac{a = upd^{RA}(x, m, n) \quad w \in E.OW(t, x) \setminus E.covered}{wrval(w) = m \quad rf' = rf \cup \{(w, e)\} \quad mo' = mo[w, e]} \\
\hline
((D, sb), rf, mo) \xrightarrow{e} ((X, sb) + e, rf', mo')
\end{array}$$

■ **Figure 14** Event semantics assuming  $E = ((X, sb), rf, mo)$ ,  $e = (g, a, t)$  and  $g \notin tags(X)$

971      The set of *encountered writes* are the writes that thread  $t$  is aware of (either directly or  
 972 indirectly) in state  $E = ((X, sb), rf, mo)$ , and are given by:

$$\begin{array}{c}
973 \quad E.EW = \{w \in W \cap D \mid \exists e \in D. \text{tid}(e) = t \wedge \\
974 \quad \quad \quad (w, e) \in \text{eco}^?; \text{hb}^?\} \cup I_E
\end{array}$$

975 where  $R^?$  is the reflexive closure of relation  $R$ . Thus, for each  $w \in E.EW$ , there must exist  
 976 an event  $e$  of thread  $t$  such that  $w$  is either *eco*- or *hb*- or *eco;hb*-prior to  $e$ .

977      From these, we determine the *observable writes*, which are the writes that thread  $t$  can  
 978 observe in its next read. These are defined as:

$$\begin{array}{c}
979 \quad E.OW(t, x) = \{w \in W \cap E.X \mid \text{loc}(w) = x \wedge \forall w' \in E.EW(t) \wedge (w, w') \notin E.mo\} \\
980
\end{array}$$

981 Thus, observable writes are not succeeded by any encountered write in modification order,  
 982 i.e., the thread has not seen another write overwriting the value being read.

983      Finally, to guarantee *atomicity* of the update events, there cannot be any write operations  
 984 (in modification order) between the write that an update reads from and the write of the  
 985 update itself. We therefore define the set of *covered writes* as follows:

$$\begin{array}{c}
986 \quad E.covered = \{w \in W \cap E.X \mid \exists u \in U. (w, u) \in rf\} \\
987
\end{array}$$

988      The transition relation of the Explicit semantics is given in Figure 14.

989 ► **Lemma 15** (Invariants of the Explicit Semantics). *If  $E = (X, sb, rf, mo)$  and  $E$  is reachable  
 990 from an initial state via a sequence of transitions of the Explicit semantics, then*

991 ■  *$mo$  totally orders the writes to each variable  $x$ . That is, for all  $w, w' \in X \cap W$ , such that  
 992  $\text{loc}(w) = \text{loc}(w')$ .*

$$\begin{array}{c}
993 \quad (w, w') \in mo \vee (w', w) \in mo \\
994
\end{array} \tag{11}$$

995 ■ *For each variable  $x$ , there is at least one write. That is, for each  $x$*

$$\begin{array}{c}
996 \quad \exists w \in X \cap W. \text{loc}(w) = x \\
997
\end{array} \tag{12}$$

998 **Proof.** These are simple inductive invariants of the semantics. ◀

## A.2 Soundness of the View Semantics

We turn now to the soundness of the View semantics. Note that in the semantics as presented in the body of this paper, writes are recorded in the state as a pair  $(w, q) \in W \times Q$ . For uniformity with the Explicit state semantics, in this proof we record writes in the state as a triple  $(w, q, t) \in W \times Q \times T$ . The transition rule for writes for the View semantics now becomes

$$\text{WRITE} \frac{a \in \{wr(x, n), wr^R(x, n)\} \quad (w, q) \in \sigma.OW(t, x) \setminus \sigma.covered \quad fresh(q, q') \quad \begin{array}{l} writes' = \sigma.writes \cup \{(a, q', t)\} \\ tview'_t = \sigma.tview_t[x := (a, q', t)] \end{array}}{\sigma \xrightarrow{a}_t \sigma[tview_t := tview'_t, mview_{(a, q')} := tview'_t, writes := writes']}$$

The other rules are changed similarly. This transformation has no effect on the observable behaviour of the semantics. But now, the set of writes  $V.writes$  in some View state now has the same type (or structure) as the set of events in some Explicit state  $E.X$ , which will prove to be convenient later.

In this section, we prove the following theorem, which states that every behaviour of the View semantics is also a behaviour of the Explicit semantics.

► **Theorem 16.** *For every sequence of steps of the View semantics*

$$V_0 \xrightarrow{a_1} V_1 \xrightarrow{a_2} \dots \xrightarrow{a_n} V_n$$

such that  $V_0$  is an initial state of the View semantics, there is a sequence of steps of the Explicit semantics

$$V_0 \xrightarrow{a_1} V_1 \xrightarrow{a_2} \dots \xrightarrow{a_n} V_n$$

such that  $E_0$  is an initial state of the Explicit semantics, having the same sequence of actions.

We prove this using the method of *simulation relations* (see e.g., [24]).

First some preliminary definitions.

► **Definition 17.** *Given an Explicit state  $E = (X, sb, rf, mo)$*

1. (*Events of a thread.*) Let  $E.X_t = \{e \in X \mid tid(e) = t\}$
  2. (*mo-maximal.*) For  $S \subseteq X$ , let  $E.\mathbf{max}_{mo}(S) = \{w \in S \cap W \mid \forall w' \in X. (w, w') \notin mo\}$
  3. (*Causal closure.*) For  $S \subseteq X$ , let  $E.\mathbf{cclose}(S) = \{e \mid \exists e' \in S. (e, e') \in eco^? \circ hb^?\} \cup I_E$
  4. (*View modification order.*) Let  $V.mo = \{(w, w') \mid w, w' \in V.writes \wedge tst(w) < tst(w')\}$
- where  $P.eco$  and  $P.hb$  are defined for an C11 RAR execution as usual.

We sometimes omit the Explicit state from these auxiliary variables, when the state in question is clear from context. For example,  $E.\mathbf{max}_{mo}$  sometimes becomes  $\mathbf{max}_{mo}$ .

► **Lemma 18.** *For any Explicit state  $E = (X, sb, rf, mo)$*

$$EW(t) = \mathbf{cclose}(X_t) \cap W \tag{13}$$

and for any set  $S \subseteq X$ ,

1.

$$\mathbf{max}_{mo}(S) \subseteq X \tag{14}$$

2.

$$\mathbf{max}_{mo}(S) \subseteq W \tag{15}$$



1038 3. for each variable  $x$  there is precisely one write  $w \in \mathbf{max}_{\mathbf{mo}}(S)$  such that  $\text{loc}(w) = x$ .

1039 **Proof.** We prove each property in turn.

1040 ■ (13) This follows immediately from the definitions of  $EW$  and  $\mathbf{cclose}$ :

$$\begin{aligned}
 1041 \quad w \in EW(t) &\iff w \in W \wedge \exists e' \in X_t.(e, e') \in \mathbf{eco}^?; \mathbf{hb} \\
 1042 &\iff w \in W \wedge w \in \mathbf{cclose}(X_t) \\
 1043 &\iff w \in \mathbf{cclose}(X_t) \cap W \\
 1044
 \end{aligned}$$

1045 ■ (14) By the definition of  $\mathbf{max}_{\mathbf{mo}}$ , if  $w \in \mathbf{max}_{\mathbf{mo}}(S)$  then  $w \in S \subseteq X$

1046 ■ (15) By the definition of  $\mathbf{max}_{\mathbf{mo}}$ , if  $w \in \mathbf{max}_{\mathbf{mo}}(S)$  then  $w \in W$ .

1047 ■ (3) By Property 11 of the explicit semantics,  $E.\mathbf{mo}$  totally orders the writes to each  $x$ .  
 1048 Thus, for any distinct writes  $v, w$  such that  $\text{loc}(v) = \text{loc}(w) = x$ , either  $(v, w) \in E.\mathbf{mo}$  or  
 1049  $(w, v) \in E.\mathbf{mo}$ . In each case, the  $\mathbf{mo}$ -earlier of the two writes cannot be  $\mathbf{mo}$ -maximal, and  
 1050 therefore  $v$  and  $w$  cannot both be in  $E.\mathbf{max}_{\mathbf{mo}}(S)$ . This ensures uniqueness. The fact  
 1051 that a write to  $x$  exists is immediate from Property 12 of the Explicit semantics.

1052 ◀

1053 Property 3 shows that  $\mathbf{max}_{\mathbf{mo}}(X)$  defines a function from variables to writes. We denote by  
 1054  $\mathbf{max}_{\mathbf{mo}}(X, x)$  the unique write in  $\mathbf{max}_{\mathbf{mo}}(X)$  such that  $\text{loc}(w) = x$ .

1055 For any Explicit state  $E$ , and any  $S \subseteq E.X \cap W$ , we say that  $S$  is *complete* if for every  
 1056 location  $x$  there is some  $w \in S$  with  $x = \text{loc}(w)$ . Note that if  $S$  is complete then  $\mathbf{max}_{\mathbf{mo}}(S, x)$   
 1057 is defined.

1058 ► **Lemma 19.** For any Explicit state  $E$ , and any  $S \subseteq E.X \cap W$ ,  $E.\mathbf{cclose}(S)$  is complete.

1059 **Proof.**  $I_E \subseteq E.\mathbf{cclose}(S)$ , and  $I_E$  is clearly complete. ◀

1060 ► **Lemma 20.** For any Explicit state  $E$ , and any complete  $S, S' \subseteq E.X$  where  $S$  is complete,  
 1061 if there exists a  $w \in S'$  such that  $\text{loc}(w) = x$ ,

$$\begin{aligned}
 1062 \quad &E.\mathbf{max}_{\mathbf{mo}}(S \cup S', x) \\
 1063 \quad &= \begin{cases} E.\mathbf{max}_{\mathbf{mo}}(S, x) & \text{if } (E.\mathbf{max}_{\mathbf{mo}}(S', x), E.\mathbf{max}_{\mathbf{mo}}(S, x)) \in E.\mathbf{mo} \\ E.\mathbf{max}_{\mathbf{mo}}(S', x) & \text{otherwise} \end{cases} \quad (16) \\
 1064
 \end{aligned}$$

1065 otherwise

$$\begin{aligned}
 1066 \quad &E.\mathbf{max}_{\mathbf{mo}}(S \cup S', x) = E.\mathbf{max}_{\mathbf{mo}}(S, x) \quad (17) \\
 1067
 \end{aligned}$$

1068 **Proof.** If there exists a  $w \in S'$  such that  $\text{loc}(w) = x$  then both  $E.\mathbf{max}_{\mathbf{mo}}(S', x)$  and  
 1069  $E.\mathbf{max}_{\mathbf{mo}}(S, x)$  are defined and  $E.\mathbf{max}_{\mathbf{mo}}(S \cup S', x)$  is the maximum of these.

1070 If there is no such write then

$$1071 \quad (S \cup S') \cap \{w \mid \text{loc}(w) = x\} = S \cap \{w \mid \text{loc}(w) = x\}$$

1072 and thus  $E.\mathbf{max}_{\mathbf{mo}}(S \cup S', x) = E.\mathbf{max}_{\mathbf{mo}}(S, x)$  as required. ◀

1073 Note that for every step of the Explicit semantics  $E \xRightarrow{e} E'$  there is some unique write  
 1074 that the event  $e$  interacts with, for example the write that  $e$  reads from if  $e$  is a read or  
 1075 RMW. The write  $w$  mentioned in each of the Explicit semantics transition rules of Figure 14.  
 1076 In what follows we exhibit this write as a label on the transition relation. Thus we write  
 1077  $E \xRightarrow{w, e} E'$  iff  $E \xRightarrow{e} E'$  and  $w$  is the write that  $e$  interacts with.

1078 ► **Definition 21.** Given an Explicit state  $E$ , we say that a thread  $t$  is before a write  $w$  at  
 1079 location  $x$ , denoted  $t \preceq_x w$  if

$$1080 \quad (E.\mathbf{max}_{\text{mo}}(E.EW(t), x), E.\mathbf{max}_{\text{mo}}(E.\mathbf{cclose}(\{w\}), x)) \in E.\text{mo}$$

1081 Likewise, we say that a write  $w$  is before a write  $t$  at location  $x$ , denoted  $w \preceq_x t$  if

$$1082 \quad (E.\mathbf{max}_{\text{mo}}(E.\mathbf{cclose}(\{w\}), x), E.\mathbf{max}_{\text{mo}}(E.EW(t), x)) \in E.\text{mo}$$

1083 ► **Lemma 22** (Max Encountered Writes). For any Explicit transition  $E \xRightarrow{w,e} E'$  where  $w$  and  
 1084  $e$  synchronise, and for all  $x$ , if  $e$  is an (acquiring) read then

$$1085 \quad E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x) = \begin{cases} E.\mathbf{max}_{\text{mo}}(E.EW(t), x) & \text{if } w \preceq_x t \\ E.\mathbf{max}_{\text{mo}}(E.\mathbf{cclose}(\{w\}), x) & \text{otherwise} \end{cases} \quad (18)$$

1086 and if  $e$  is an update then

$$1087 \quad E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x) \quad (19)$$

$$1088 \quad = \begin{cases} e & \text{if } \text{loc}(w) = x \\ E.\mathbf{max}_{\text{mo}}(E.EW(t), x) & \text{if } \text{loc}(w) \neq x \wedge w \preceq_x t \\ E.\mathbf{max}_{\text{mo}}(E.\mathbf{cclose}(\{w\}), x) & \text{otherwise} \end{cases}$$

1089 Further, for any Explicit transition  $E \xRightarrow{w,e} E'$  where  $w$  and  $e$  do not synchronise, if  $e$  is a  
 1090 read then

$$1091 \quad E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x) = \begin{cases} w & \text{if } \text{loc}(w) = x \\ E.\mathbf{max}_{\text{mo}}(E.EW(t), x) & \text{otherwise} \end{cases} \quad (20)$$

1092 and if  $e$  is a write or update then

$$1093 \quad E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x) = \begin{cases} e & \text{if } \text{loc}(w) = x \\ E.\mathbf{max}_{\text{mo}}(E.EW(t), x) & \text{otherwise} \end{cases} \quad (21)$$

1094 **Proof.** We prove each in turn.

1095 (Equation 18) Consider

$$1096 \quad E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x)$$

$$1097 \quad = \quad (\text{trans-rel})$$

$$1098 \quad E.\mathbf{max}_{\text{mo}}(E'.EW(t), x)$$

$$1099 \quad = \quad (\text{trans-rel and defn of } \mathbf{cclose})$$

$$1100 \quad E.\mathbf{max}_{\text{mo}}(E.EW(t) \cup E.\mathbf{cclose}(\{w\}), x)$$

$$1101 \quad = \quad (\text{see below})$$

$$1102 \quad \begin{cases} E.\mathbf{max}_{\text{mo}}(E.EW(t), x) & \text{if } w \preceq_x t \\ E.\mathbf{max}_{\text{mo}}(E.\mathbf{cclose}(\{w\}), x) & \text{otherwise} \end{cases}$$

1103 This last step is a consequence of Lemma 20 and the fact that both  $E.EW(t)$  and  $E.\mathbf{cclose}(\{w\})$   
 1104 are complete.

(Equation 19) We first consider the case when  $x = \text{loc}(e)$ . In this case

$$E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x) = e$$

1110 as required. If  $x \neq \text{loc}(e)$ , then

$$\begin{aligned}
1111 & E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x) \\
1112 & = \quad (\text{trans-rel}) \\
1113 & E.\mathbf{max}_{\text{mo}}(E'.EW(t), x) \\
1114 & = \quad (\text{trans-rel and defn of } \mathbf{cclose}) \\
1115 & E.\mathbf{max}_{\text{mo}}(E.EW(t) \cup E.\mathbf{cclose}(\{w\}) \cup \{e\}, x) \\
1116 & = \quad (x \neq \text{loc}(e) \text{ and Lemma 20}) \\
1117 & E.\mathbf{max}_{\text{mo}}(E.EW(t) \cup E.\mathbf{cclose}(\{w\}), x) \\
1118 & = \quad (\text{see below}) \\
1119 & \begin{cases} E.\mathbf{max}_{\text{mo}}(E.EW(t), x) & \text{if } w \preceq_x t \\ E.\mathbf{max}_{\text{mo}}(E.\mathbf{cclose}(\{w\}), x) & \text{otherwise} \end{cases} \\
1120 &
\end{aligned}$$

1121 Agin, this last step is a consequence of Lemma 20 and the fact that both  $E.EW(t)$  and  
 1122  $E.\mathbf{cclose}(\{w\})$  are complete.

(Equation 20) We first consider the case when  $x = \text{loc}(e)$ . In this case

$$E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x) = e$$

1123 as required. If  $x \neq \text{loc}(e)$ , then

$$\begin{aligned}
1124 & E'.\mathbf{max}_{\text{mo}}(E'.EW(t), x) \\
1125 & = \quad (\text{trans-rel}) \\
1126 & E.\mathbf{max}_{\text{mo}}(E'.EW(t), x) \\
1127 & = \quad (\text{trans-rel and defn of } \mathbf{cclose}) \\
1128 & E.\mathbf{max}_{\text{mo}}(E.EW(t) \cup \{w\} \cup \{e\}, x) \\
1129 & = \quad (\text{because } x \neq \text{loc}(e) \text{ and Lemma 20}) \\
1130 & E.\mathbf{max}_{\text{mo}}(E.EW(t), x) \\
1131 &
\end{aligned}$$

1132 (Equation 21) The proof here is identical to that for Equation 20. ◀

1133 We define a simulation relation  $R$  between a View state  $V = (\text{writes}, \text{tview}, \text{mview}, \text{covered})$   
 1134 and an Explicit state  $E = (X, \text{sb}, \text{rf}, \text{mo})$  to be the conjunction of the following properties.

1135 ■ Both executions contain the same set of writes:

$$1136 \quad V.\text{writes} = E.X \cap W \tag{22}$$

1138 ■ For all threads  $t$  and variables  $x$ ,

$$1139 \quad V.\text{tview}(t, x) = E.\mathbf{max}_{\text{mo}}(E.EW(t), x) \tag{23}$$

1141 ■ For all  $w \in \text{writes}$  and variables  $x$ ,

$$1142 \quad V.\text{mview}(w, x) = E.\mathbf{max}_{\text{mo}}(E.\mathbf{cclose}(\{w\}), x) \tag{24}$$

1144 ■ Both executions agree on the order of writes.

$$1145 \quad V.\text{mo} = E.\text{mo} \tag{25}$$

1147 recall that  $V.\text{mo} = \{(w, w') \mid w, w' \in V.\text{writes} \wedge \text{tst}(w) < \text{tst}(w')\}$ .

1148 ■ The executions agree on the set of covered writes:

$$1149 \quad V.covered = E.covered \quad (26)$$

1151 Our first lemma relates the viewfront and the observable writes of  $R$ -related states.

1152 ► **Lemma 23.** *For any View state  $V$  and Explicit state  $E$  such that  $(V, E) \in R$ , and for all*  
 1153 *threads  $t$  and locations  $x$ ,  $V.OW(t, x) = E.OW(t, x)$ .*

1154 **Proof.** First, observe that for all  $w$  such that  $loc(w) = x$

$$1155 \quad tst(V.tview_t(x)) \leq tst(w) \iff (E.\mathbf{max}_{\mathbf{mo}}(E.EW(t), x), w) \in E.\mathbf{mo} \quad (27)$$

1157 But this is an immediate consequence of the fact that  $V.tview(t, x) = E.\mathbf{max}_{\mathbf{mo}}(E.EW(t), x)$   
 1158 and that  $E.\mathbf{mo} = V.\mathbf{mo}$  (both of these are consequences of  $R$ ).

1159 But now, because  $E.\mathbf{max}_{\mathbf{mo}}(E.EW(t), x)$  is  $\mathbf{mo}$ -maximal among the set of encountered  
 1160 writes, we have

$$1161 \quad (E.\mathbf{max}_{\mathbf{mo}}(E.EW(t), x), w) \in E.\mathbf{mo} \iff \forall w' \in E.EW(t), x. (w, w') \notin E.\mathbf{mo} \quad (28)$$

1163 Now consider

$$\begin{aligned} 1164 \quad & w \in V.OW(t, x) \\ 1165 \quad & \iff \text{(definition)} \\ 1166 \quad & w \in V.writes \wedge loc(a) = x \wedge tst(V.tview_t(x)) \leq tst(w) \\ 1167 \quad & \iff \text{(by } R) \\ 1168 \quad & w \in E.X \cap W \wedge loc(a) = x \wedge tst(V.tview_t(x)) \leq tst(w) \\ 1169 \quad & \iff (28) \\ 1170 \quad & w \in E.X \cap W \wedge loc(a) = x \wedge \forall w' \in E.EW(t), x. (w, w') \notin E.\mathbf{mo} \\ 1171 \quad & \iff w \in E.OW(t, x) \end{aligned}$$

1173 as required. ◀

1174 In the preservation proof, we use the following stability properties:

1175 ► **Lemma 24.** *For all  $E \xrightarrow{w, e} E'$ , every location  $x$ , and every thread  $t' \neq tid(e)$*

$$1176 \quad E'.\mathbf{max}_{\mathbf{mo}}(E'.EW(t'), x) = E.\mathbf{max}_{\mathbf{mo}}(E.EW(t'), x) \quad (29)$$

1178 and for every write  $w' \neq e$

$$1179 \quad E'.\mathbf{max}_{\mathbf{mo}}(E'.\mathbf{cclose}(\{w'\}), x) = E.\mathbf{max}_{\mathbf{mo}}(E.\mathbf{cclose}(\{w'\}), x) \quad (30)$$

1181 Further, for all  $V \xrightarrow{w, e} V'$ , every location  $x$ , and every thread  $t' \neq tid(e)$

$$1182 \quad V'.tview(t', x) = V.tview(t', x) \quad (31)$$

1184 and for every write  $w' \neq e$

$$1185 \quad V'.mview(w', x) = V.mview(w', x) \quad (32)$$

1187 ► **Lemma 25** (Mod-order agreement). *For any View state*

$$1188 \quad V = (\text{writes}, \text{tview}, \text{mview}, \text{covered})$$

1189 *and Explicit state  $E = (X, \text{sb}, \text{rf}, \text{mo})$  such that  $(V, E) \in R$ , and every thread  $t$ , write  $w$  and*  
 1190 *variable  $x$*

$$1191 \quad \text{tst}(V.\text{tview}(t)(x)) < \text{tst}(V.\text{mview}(w)(x)) \iff t \preceq_x w$$

1193 **Proof.** We reason thusly

$$1194 \quad \text{tst}(V.\text{tview}(t)(x)) < \text{tst}(V.\text{mview}(w)(x)) \iff$$

$$1195 \quad \text{tst}(E.\text{max}_{\text{mo}}(E.EW(t), x)) < \text{tst}(E.\text{max}_{\text{mo}}(E.EW(t), x)) \iff \text{by 23}$$

$$1196 \quad t \preceq_x w \quad \text{by 25}$$

1198 ◀

1199 ► **Lemma 26** (View Mod Order). *For any  $V \xrightarrow{w,e}_t V'$  where  $e$  is a write or update,  $V'.\text{mo} =$*   
 1200  *$V.\text{mo}[w, (e, q, t)]$  where  $q$  is the fresh rational used to tag the operation in the View transition*  
 1201 *relation.*

1202 **Proof.** The new write is added into **mo** immediately after the write  $w$  and before all  
 1203 subsequent writes to the same variable. ◀

1204 ► **Lemma 27** ( $R$ -preservation). *For any View state*

$$1205 \quad V = (\text{writes}, \text{tview}, \text{mview}, \text{covered})$$

1206 *and Explicit state  $E = (X, \text{sb}, \text{rf}, \text{mo})$  such that  $(V, E) \in R$ , and every View state  $V' =$*   
 1207  *$(\text{writes}', \text{tview}', \text{mview}', \text{covered}')$  such that  $V \xrightarrow{w,e}_t V'$ , we have  $E \xrightarrow{w,e}_t E'$  for some  $E'$  and*  
 1208  *$(V', E') \in R$ .*

1209 **Proof.** We proceed by cases on the type of the operation  $e$  and whether or not the operation  
 1210 is synchronising. In what follows, let  $t = \text{tid}(e)$ .

1211 **Case 1.**  $e$  is of the form  $rd(x, n)$ . Let  $E = (X', \text{sb}', \text{rf}', \text{mo}')$  where

$$1212 \quad (X', \text{sb}') = (X, \text{sb}) + e$$

$$1213 \quad \text{rf}' = \text{rf} \cup \{(w, e)\}$$

$$1214 \quad \text{mo}' = \text{mo}$$

1216 The precondition of the View transition ensures that  $w \in V.OW(t, x)$ , and thus by Lemma  
 1217 23, we have  $w \in E.OW(t)$ . Thus, we have  $E \xrightarrow{w,e}_t E'$ . It remains to show that  $(V', E') \in R$ ,  
 1218 which we do by considering each of the equations in the definition of  $R$  in turn.

■

$$1219 \quad V'.\text{writes} = V.\text{writes} \quad \text{by View trans-rel}$$

$$1220 \quad = E.X \cap W \quad \text{by } R$$

$$1221 \quad = E'.X \cap W \quad \text{by Explicit trans-rel}$$

1223 ■ We need to show that  $V'.\text{tview}(t, x) = E'.\text{max}_{\text{mo}}(E'.EW(t), x)$ . For all  $t' \neq t$ , it is easy  
 1224 to see that Equations 29 and 31 ensure that this property is preserved.

1225 Several cases remain. We discuss the first two. The remaining cases follow in a similar  
 1226 way. In the first case, we assume the following

$$1227 \quad w, e \text{ synchronise} \quad (33)$$

$$1228 \quad \text{loc}(e) = x \quad (34)$$

$$1229 \quad \text{tst}(V.\text{tview}(t)(x)) < \text{tst}(V.\text{mview}(w)(x)) \quad (35)$$

$$1230 \quad (E.\text{max}_{\text{mo}}(E.EW(t), x), E.\text{max}_{\text{mo}}(E.\text{cclose}(\{w\}), x)) \in E.\text{mo} \quad (36)$$

1232 Note that by 25 the last two assumptions are equivalent. Then,

$$\begin{aligned} 1233 \quad V'.\text{tview}(t)(x) &= (V.\text{tview}(t) \otimes V.\text{mview}(w))(x) && \text{by View trans-rel} \\ 1234 \quad &= V.\text{mview}(w)(x) && \text{by hyp. and def of } \otimes \\ 1235 \quad &= E.\text{max}_{\text{mo}}(E.\text{cclose}(\{w\}), x) && \text{by } R \\ 1236 \quad &= E'.\text{max}_{\text{mo}}(E'.EW(t), x) && \text{see below} \end{aligned}$$

1238 The last step is a consequence of the first, second, and fourth assumptions together with  
 1239 Lemma 18.

1240 In the second case, we assume

$$1241 \quad w, e \text{ synchronise} \quad (37)$$

$$1242 \quad \text{loc}(e) = x \quad (38)$$

$$1243 \quad \text{tst}(V.\text{mview}(w)(x)) < \text{tst}(V.\text{tview}(t)(x)) \quad (39)$$

$$1244 \quad (E.\text{max}_{\text{mo}}(E.\text{cclose}(\{w\}), x), E.\text{max}_{\text{mo}}(E.EW(t), x)) \in E.\text{mo} \quad (40)$$

1246 (The new hypothesis differs from the previous in that we have changed the view that  
 1247 supplies the latest write.) Note that by 25 the last two assumptions are equivalent. Then,

$$\begin{aligned} 1248 \quad V'.\text{tview}(t)(x) &= (V.\text{tview}(t) \otimes V.\text{mview}(w))(x) && \text{by View trans-rel} \\ 1249 \quad &= V.\text{tview}(t)(x) && \text{by hyp. and def of } \otimes \\ 1250 \quad &= E.\text{max}_{\text{mo}}(E.EW(t), x) && \text{by } R \\ 1251 \quad &= E'.\text{max}_{\text{mo}}(E'.EW(t), x) && \text{see below} \end{aligned}$$

1253 The last step is a consequence of the first, second, and fourth assumptions together with  
 1254 Lemma 18.

1255 ■ We need to show that  $V'.\text{mview}(w, x) = E'.\text{max}_{\text{mo}}(E'.\text{cclose}(\{w\}), x)$  for all  $w$ . But  
 1256 because  $e$  is a read we have

$$1257 \quad V'.\text{mview} = V.\text{mview}$$

$$1258 \quad E'.\text{cclose} = E.\text{cclose}$$

$$1259 \quad E'.\text{max}_{\text{mo}} = E.\text{max}_{\text{mo}}$$

1261 so the preservation result follows immediately.

$$\begin{aligned} 1262 \quad V'.\text{mo} &= V.\text{mo} && \text{by View trans-rel} \\ 1263 \quad &= E.\text{mo} && \text{by } R \\ 1264 \quad &= E'.\text{mo} && \text{by Explicit trans-rel} \end{aligned}$$

$$\begin{array}{ll}
1266 & V'.covered = V.covered \quad \text{by View trans-rel} \\
1267 & = E.covered \quad \text{by } R \\
1268 & = E'.covered \quad \text{by Explicit trans-rel} \\
1269 &
\end{array}$$

1270 **Case 2.**  $e$  is of the form  $wr(x, n)$ . Let  $E = (X', sb', rf', mo')$  where

$$\begin{array}{ll}
1271 & (X', sb') = (X, sb) + e \\
1272 & rf' = rf \\
1273 & mo' = mo[w, e] \\
1274 &
\end{array}$$

1275 The precondition of the View transition ensures that  $w \in V.OW(t, x)$ , and thus by Lemma  
1276 23, we have  $w \in E.OW(t)$ . Thus, we have  $E \xrightarrow{w, e} E'$ . It remains to show that  $(V', E') \in R$ ,  
1277 which we do by considering each of the equations in the definition of  $R$  in turn.

$$\begin{array}{ll}
1278 & V'.writes = V.writes \cup \{e\} \quad \text{by View trans-rel} \\
1279 & = (E.X \cap W) \cup \{e\} \quad \text{by } R \\
1280 & = E'.X \cap W \quad \text{by Explicit trans-rel} \\
1281 &
\end{array}$$

■ If  $loc(e) \neq x$  then the thread view and  $mo$  restricted to  $x$  are unchanged. If  $loc(e) = x$ , then  $E'.\max_{mo}(E'.EW(t), x) = e$  and  $V'.tview(t, x) = e$  so

$$E'.\max_{mo}(E'.EW(t), x) = V'.tview(t, x) = e$$

1282 as required.

■ If  $loc(e) \neq x$  then the thread view and  $mo$  restricted to  $x$  are unchanged. Assume  $loc(e) = x$ . For all  $w' \in V'.writes$  where  $w' \neq e$ , then

$$E'.\max_{mo}(E'.cclose(w), x) = E.\max_{mo}(E.cclose(w), x)$$

so the property is preserved. In the final case,

$$E'.\max_{mo}(E'.cclose(e), x) = E.\max_{mo}(E.EW(t) \cup \{e\}) = e$$

1283 but  $V'.mview(e, x) = e$  as required.

$$\begin{array}{ll}
1284 & V'.mo = V.mo[w, e] \quad \text{by Lemma 26} \\
1285 & = E.mo[w, e] \quad \text{by } R \\
1286 & = E'.mo \quad \text{by Explicit trans-rel} \\
1287 &
\end{array}$$

$$\begin{array}{ll}
1288 & V'.covered = V.covered \quad \text{by View trans-rel} \\
1289 & = E.covered \quad \text{by } R \\
1290 & = E'.covered \quad \text{by Explicit trans-rel} \\
1291 &
\end{array}$$



1292 **Case 3.**  $e$  is of the form  $upd^{RA}(x, m, n)$ . Let  $E = (X', sb', rf', mo')$  where

$$1293 (X', sb') = (X, sb) + e$$

$$1294 rf' = rf \cup \{(w, e)\}$$

$$1295 mo' = mo[w, e]$$

1297 The precondition of the View transition ensures that  $w \in V.OW(t, x)$ , and thus by Lemma  
 1298 23, we have  $w \in E.OW(t)$ . Thus, we have  $E \xrightarrow{w, e} E'$ . It remains to show that  $(V', E') \in R$ ,  
 1299 which we do by considering each of the equations in the definition of  $R$  in turn.

1300 The proof that  $(V', E') \in R$  is essentially a combination of the proof for reads and writes.

$$\begin{aligned} 1301 V'.writes &= V.writes \cup \{e\} && \text{by View trans-rel} \\ 1302 &= (E.X \cap W) \cup \{e\} && \text{by } R \\ 1303 &= E'.X \cap W && \text{by Explicit trans-rel} \end{aligned}$$

1305 ■ The proof here is the same as that for reads.

1306 ■ The proof here is the same as that for writes.

$$\begin{aligned} 1307 V'.mo &= V.mo[w, e] && \text{by Lemma 26} \\ 1308 &= E.mo[w, e] && \text{by } R \\ 1309 &= E'.mo && \text{by Explicit trans-rel} \end{aligned}$$

$$\begin{aligned} 1311 V'.covered &= V.covered \cup \{w\} && \text{by View trans-rel} \\ 1312 &= E.covered \cup \{w\} && \text{by } R \\ 1313 &= E'.covered && \text{by Explicit trans-rel} \end{aligned}$$

1315

1316 Finally, we must prove that for every initial View state, there is an  $R$ -related initial  
 1317 Explicit state.

1318 ► **Lemma 28.** *For every initial state of the View semantics  $V_0$ , there is an initial state of*  
 1319 *the Explicit semantics  $E_0$  such that  $(V_0, E_0) \in R$*

1320 **Proof.** Let  $V_0 = (writes, tview, mview, covered)$ . We let  $E_0 = (X, sb, rf, mo)$ , where <sup>6</sup>

$$1321 X = writes \tag{41}$$

$$1322 sb = \emptyset \tag{42}$$

$$1323 rf = \emptyset \tag{43}$$

$$1324 mo = \emptyset \tag{44}$$

1326 We prove each property of the simulation relation in turn.

1327 ■ (22) This is immediate.

<sup>6</sup> Recall that in our setting  $writes$  and  $X$  have the same type.

## 0:38      Owicki-Gries Reasoning for C11 RAR

- 1328   ■   (23) For all threads  $t$  and variables  $x$ ,  $E.\mathbf{max}_{\text{mo}}(E.EW(t), x)$  is the initialising write to  $x$ ,  
1329       which is also the value of  $V.tview(t, x)$ .
- 1330   ■   (24) Similarly to 23, for all writes  $w$  and variables  $x$ ,  $E.\mathbf{cclose}(w)$  and  $V.mview(w)$  is the  
1331       initialising write to  $x$ .
- 1332   ■   (25) This is immediate.
- 1333   ■   (26) This is immediate.
- 1334   